

# Advanced Kubernetes

## Lab 3 – Scheduler

Per the k8s reference documentation, the scheduler is described as follows:

The Kubernetes scheduler is a policy-rich, topology-aware, workload-specific function that significantly impacts availability, performance, and capacity. The scheduler needs to take into account individual and collective resource requirements, quality of service requirements, hardware/software/policy constraints, affinity and anti-affinity specifications, data locality, inter-workload interference, deadlines, and so on. Workload-specific requirements will be exposed through the API as necessary.

<http://kubernetes.io/docs/admin/kube-scheduler/>

In this lab we will see how the scheduler affects placement of pods.

### 1. Deploy a two node cluster

In order to demonstrate multi-node cluster operations and pod scheduling we will set up a second node called *nodeb* (The master node was configured in lab 1 and was called *nodea*).

### 2. Update IPs and hostnames

Set the host name for the new VM to *nodeb*:

```
laptop$ ssh -i k8s-adv-student.pem ubuntu@<external-ip>  
  
...  
  
ubuntu@nodeb:~$ sudo hostnamectl set-hostname nodeb
```

```
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ hostname
```

```
nodeb
```

```
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ cat /etc/hostname
```

```
nodeb
```

```
ubuntu@nodeb:~$
```

Now discover your IP address (typically eth0 or ens33):

```
ubuntu@nodeb:~$ ip a show eth0
```

```
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc mq state UP group default qlen 1000
    link/ether 02:d8:c0:66:a6:b8 brd ff:ff:ff:ff:ff:ff
    inet 172.31.30.148/20 brd 172.31.31.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::d8:c0ff:fe66:a6b8/64 scope link
        valid_lft forever preferred_lft forever
ubuntu@nodeb:~$
```

Add your IP address and host name to `/etc/hosts`, also add *nodea*'s information and remove any references to the ubuntu hostname. In a more sophisticated setting, DNS could be used to perform hostname lookups.

```
ubuntu@nodeb:~$ sudo vim /etc/hosts
```

```
ubuntu@nodeb:~$ cat /etc/hosts
```

```
127.0.0.1    localhost
172.31.30.148 nodeb
172.31.28.198 nodea
```

```
# The following lines are desirable for IPv6 capable hosts
::1 ip6-localhost ip6-loopback
fe00::0 ip6-localnet
ff00::0 ip6-mcastprefix
ff02::1 ip6-allnodes
ff02::2 ip6-allrouters
ff02::3 ip6-allhosts
```

You may need to exit the current shell and open a new shell for your prompt (PS1) to update to the new hostname.

Now go back to *nodea*, add *nodeb*'s IP information to the `/etc/hosts` file. Depending on the hypervisor and technique used your IPs may differ.

Finally, verify that you can reach the internet and both nodes by name with ping from both VMs:

```
ubuntu@nodeb:~$ ping -c 2 yahoo.com

PING yahoo.com (98.137.246.7) 56(84) bytes of data.
64 bytes from media-router-fp1.prod1.media.vip.gq1.yahoo.com (98.137.246.7): icmp_seq=1 ttl=43 time=11.7 ms
64 bytes from media-router-fp1.prod1.media.vip.gq1.yahoo.com (98.137.246.7): icmp_seq=2 ttl=43 time=11.8 ms

--- yahoo.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 11.741/11.776/11.811/0.035 ms
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ ping -c 2 nodea

PING nodea (172.31.28.198) 56(84) bytes of data.
64 bytes from nodea (172.31.28.198): icmp_seq=1 ttl=64 time=0.370 ms
64 bytes from nodea (172.31.28.198): icmp_seq=2 ttl=64 time=0.391 ms

--- nodea ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1000ms
rtt min/avg/max/mdev = 0.370/0.380/0.391/0.022 ms
ubuntu@nodeb:~$
```

```
ubuntu@nodea:~$ ping -c 2 nodeb

PING nodeb (172.31.30.148) 56(84) bytes of data:
64 bytes from nodeb (172.31.30.148): icmp_seq=1 ttl=64 time=0.451 ms
64 bytes from nodeb (172.31.30.148): icmp_seq=2 ttl=64 time=0.410 ms

--- nodeb ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 999ms
rtt min/avg/max/mdev = 0.410/0.430/0.451/0.029 ms
ubuntu@nodea:~$
```

If you can not resolve public DNS names or reach the internet, debug your connectivity before continuing.

### 3. Install Docker

Every k8s node will need Docker installed. We have already installed Docker on *nodea*, now do the same for *nodeb*. We will use a short cut script supplied by docker:

Note: if you get errors regarding dpkg your system is probably updating, wait a few minutes and try again.

```
ubuntu@nodeb:~$ sudo apt-get update

ubuntu@nodeb:~$ sudo apt-get -y install apt-transport-https ca-certificates curl

ubuntu@nodeb:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -

ubuntu@nodeb:~$ sudo add-apt-repository \
"deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"

ubuntu@nodeb:~$ sudo apt-get update

ubuntu@nodeb:~$ sudo apt-get -y install docker-ce

ubuntu@nodeb:~$ sudo usermod -aG docker ubuntu

ubuntu@nodeb:~$ exit
```

```
laptop$
```

- Does this node need to have the same version of Docker as other nodes?

The answer is no; only the kubelet on that node talks to Docker so in theory every node could have a different version of Docker. In practice it is easier to manage and debug a cluster with the same version of Docker everywhere. Some upgrade Docker versions progressively (e.g. 10% of the nodes per day) to limit the impact of latent defects or incompatibilities.

## 4. Verify Docker operation

When the system comes back up login and check the version of all parts of the Docker platform with the `docker version` subcommand:

```
laptop$ ssh -i k8s-adv-student.pem ubuntu@<external-ip>
```

```
...
```

```
ubuntu@nodeb:~$ docker version
```

```
Client:
```

```
Version:      18.09.4
API version:  1.39
Go version:   go1.10.8
Git commit:   d14af54
Built:        Wed Mar 27 18:34:51 2019
OS/Arch:      linux/amd64
Experimental: false
```

```
Server: Docker Engine - Community
```

```
Engine:
```

```
Version:      18.09.4
API version:  1.39 (minimum version 1.12)
Go version:   go1.10.8
Git commit:   d14af54
Built:        Wed Mar 27 18:01:48 2019
OS/Arch:      linux/amd64
Experimental: false
ubuntu@nodeb:~$
```

## 5. Pod placement without the scheduler

To begin, we need to restart all of the previously configured parts of the k8s cluster (etcd, kube-apiserver, kubelet).

On nodea

Stop all Kubernetes services and etcd (^C them as needed).

Clear the etcd and kubelet state caches, along with Docker containers:

```
ubuntu@nodea:~$ rm -Rf ~/default.etcd/
```

```
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ sudo rm -Rf /var/lib/kubelet/
```

```
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ docker container rm $(docker container stop $(docker container ls -qa))
```

```
...
```

```
ubuntu@nodea:~$
```

Start a fresh etcd:

```
ubuntu@nodea:~$ etcd
```

```
2019-03-30 07:44:12.491555 I | etcdmain: etcd Version: 3.3.10
2019-03-30 07:44:12.491597 I | etcdmain: Git SHA: 27fc7e2
2019-03-30 07:44:12.491612 I | etcdmain: Go Version: go1.10.4
2019-03-30 07:44:12.491620 I | etcdmain: Go OS/Arch: linux/amd64
2019-03-30 07:44:12.491629 I | etcdmain: setting maximum number of CPUs to 2, total number of available CPUs is 2
2019-03-30 07:44:12.491640 W | etcdmain: no data-dir provided, using default data-dir ./default.etcd
2019-03-30 07:44:12.491901 I | embed: listening for peers on http://localhost:2380
```

```

2019-03-30 07:44:12.491953 I | embed: listening for client requests on localhost:2379
2019-03-30 07:44:12.495983 I | etcdserver: name = default
2019-03-30 07:44:12.495999 I | etcdserver: data dir = default.etcd
2019-03-30 07:44:12.496004 I | etcdserver: member dir = default.etcd/member
2019-03-30 07:44:12.496013 I | etcdserver: heartbeat = 100ms
2019-03-30 07:44:12.496019 I | etcdserver: election = 1000ms
2019-03-30 07:44:12.496027 I | etcdserver: snapshot count = 100000
2019-03-30 07:44:12.496043 I | etcdserver: advertise client URLs = http://localhost:2379
2019-03-30 07:44:12.496052 I | etcdserver: initial advertise peer URLs = http://localhost:2380
2019-03-30 07:44:12.496064 I | etcdserver: initial cluster = default=http://localhost:2380
2019-03-30 07:44:12.499317 I | etcdserver: starting member 8e9e05c52164694d in cluster cdf818194e3a8c32
2019-03-30 07:44:12.499345 I | raft: 8e9e05c52164694d became follower at term 0
2019-03-30 07:44:12.499359 I | raft: newRaft 8e9e05c52164694d [peers: [], term: 0, commit: 0, applied: 0,
lastindex: 0, lastterm: 0]
2019-03-30 07:44:12.499368 I | raft: 8e9e05c52164694d became follower at term 1
2019-03-30 07:44:12.503658 W | auth: simple token is not cryptographically signed
2019-03-30 07:44:12.506058 I | etcdserver: starting server... [version: 3.3.10, cluster version: to_be_decided]
2019-03-30 07:44:12.506788 I | etcdserver: 8e9e05c52164694d as single-node; fast-forwarding 9 ticks (election
ticks 10)
2019-03-30 07:44:12.507107 I | etcdserver/membership: added member 8e9e05c52164694d [http://localhost:2380] to
cluster cdf818194e3a8c32
2019-03-30 07:44:13.499644 I | raft: 8e9e05c52164694d is starting a new election at term 1
2019-03-30 07:44:13.499683 I | raft: 8e9e05c52164694d became candidate at term 2
2019-03-30 07:44:13.499710 I | raft: 8e9e05c52164694d received MsgVoteResp from 8e9e05c52164694d at term 2
2019-03-30 07:44:13.499729 I | raft: 8e9e05c52164694d became leader at term 2
2019-03-30 07:44:13.499741 I | raft: raft.node: 8e9e05c52164694d elected leader 8e9e05c52164694d at term 2
2019-03-30 07:44:13.500035 I | etcdserver: published {Name:default ClientURLs:[http://localhost:2379]} to cluster
cdf818194e3a8c32
2019-03-30 07:44:13.500162 I | etcdserver: setting up the initial cluster version to 3.3
2019-03-30 07:44:13.500208 I | embed: ready to serve client requests
2019-03-30 07:44:13.500253 E | etcdmain: forgot to set Type=notify in systemd service file?
2019-03-30 07:44:13.500750 N | embed: serving insecure client requests on 127.0.0.1:2379, this is strongly
discouraged!
2019-03-30 07:44:13.500838 N | etcdserver/membership: set the initial cluster version to 3.3
2019-03-30 07:44:13.500937 I | etcdserver/api: enabled capabilities for version 3.3

```

Restart the kube-apiserver:

```

ubuntu@nodea:~$ sudo $HOME/k8s/_output/bin/kube-apiserver \
--etcd-servers=http://localhost:2379 \

```

```
--allow-privileged=true \  
--service-cluster-ip-range=10.0.0.0/16 \  
--insecure-bind-address=0.0.0.0 \  
--disable-admission-plugins=ServiceAccount
```

```
Flag --insecure-bind-address has been deprecated, This flag will be removed in a future version.  
I0330 07:45:27.185457 7701 server.go:559] external host was not specified, using 172.31.28.198  
W0330 07:45:27.185503 7701 authentication.go:415] AnonymousAuth is not allowed with the AlwaysAllow authorizer.  
Resetting AnonymousAuth to false. You should use a different authorizer  
I0330 07:45:27.185677 7701 server.go:146] Version: v1.14.0  
I0330 07:45:27.635388 7701 plugins.go:158] Loaded 7 mutating admission controller(s) successfully in the  
following order:  
NamespaceLifecycle,LimitRanger,TaintNodesByCondition,Priority,DefaultTolerationSeconds,DefaultStorageClass,Mutatin  
gAdmissionWebhook.  
I0330 07:45:27.635414 7701 plugins.go:161] Loaded 5 validating admission controller(s) successfully in the  
following order: LimitRanger,Priority,PersistentVolumeClaimResize,ValidatingAdmissionWebhook,ResourceQuota.  
E0330 07:45:27.637043 7701 prometheus.go:138] failed to register depth metric admission_quota_controller:  
duplicate metrics collector registration attempted  
E0330 07:45:27.637077 7701 prometheus.go:150] failed to register adds metric admission_quota_controller:  
duplicate metrics collector registration attempted  
E0330 07:45:27.637115 7701 prometheus.go:162] failed to register latency metric admission_quota_controller:  
duplicate metrics collector registration attempted  
E0330 07:45:27.637150 7701 prometheus.go:174] failed to register work_duration metric  
admission_quota_controller: duplicate metrics collector registration attempted  
E0330 07:45:27.637201 7701 prometheus.go:189] failed to register unfinished_work_seconds metric  
admission_quota_controller: duplicate metrics collector registration attempted  
E0330 07:45:27.637243 7701 prometheus.go:202] failed to register longest_running_processor_microseconds metric  
admission_quota_controller: duplicate metrics collector registration attempted  
I0330 07:45:27.637264 7701 plugins.go:158] Loaded 7 mutating admission controller(s) successfully in the  
following order:  
NamespaceLifecycle,LimitRanger,TaintNodesByCondition,Priority,DefaultTolerationSeconds,DefaultStorageClass,Mutatin  
gAdmissionWebhook.  
I0330 07:45:27.637275 7701 plugins.go:161] Loaded 5 validating admission controller(s) successfully in the  
following order: LimitRanger,Priority,PersistentVolumeClaimResize,ValidatingAdmissionWebhook,ResourceQuota.  
  
...
```

Restart the nodea **kubelet** :

```
ubuntu@nodea:~$ sudo $HOME/k8s/_output/bin/kubelet \  
--kubeconfig=nodea.conf \  

```



```
--config=nodea.yaml \
--allow-privileged=true \
--runtime-cgroups=/systemd/machine.slice \
--kubelet-cgroups=/systemd/machine.slice \
--pod-infra-container-image=k8s.gcr.io/pause:3.1
```

Flag `--allow-privileged` has been deprecated, will be removed in a future version

Flag `--kubelet-cgroups` has been deprecated, This parameter should be set via the config file specified by the Kubelet's `--config` flag. See <https://kubernetes.io/docs/tasks/administer-cluster/kubelet-config-file/> for more information.

Flag `--allow-privileged` has been deprecated, will be removed in a future version

Flag `--kubelet-cgroups` has been deprecated, This parameter should be set via the config file specified by the Kubelet's `--config` flag. See <https://kubernetes.io/docs/tasks/administer-cluster/kubelet-config-file/> for more information.

```
I0330 07:46:25.830948    7713 server.go:417] Version: v1.14.0
I0330 07:46:25.832351    7713 plugins.go:103] No cloud provider specified.
I0330 07:46:25.872819    7713 server.go:625] --cgroups-per-qos enabled, but --cgroup-root was not specified.
defaulting to /
I0330 07:46:25.873090    7713 container_manager_linux.go:261] container manager verified user specified cgroup-
root exists: []
I0330 07:46:25.873111    7713 container_manager_linux.go:266] Creating Container Manager object based on Node
Config: {RuntimeCgroupsName:/systemd/machine.slice SystemCgroupsName: KubeletCgroupsName:/systemd/machine.slice
ContainerRuntime:docker CgroupsPerQOS:true CgroupRoot:/ CgroupDriver:cgroupfs KubeletRootDir:/var/lib/kubelet
ProtectKernelDefaults:false NodeAllocatableConfig:{KubeReservedCgroupName: SystemReservedCgroupName:
EnforceNodeAllocatable:map[pods:{}] KubeReserved:map[] SystemReserved:map[] HardEvictionThresholds:
[{Signal:nodedefs.inodesFree Operator:LessThan Value:{Quantity:<nil> Percentage:0.05} GracePeriod:0s MinReclaim:
<nil>} {Signal:imagefs.available Operator:LessThan Value:{Quantity:<nil> Percentage:0.15} GracePeriod:0s
MinReclaim:<nil>} {Signal:memory.available Operator:LessThan Value:{Quantity:100Mi Percentage:0} GracePeriod:0s
MinReclaim:<nil>} {Signal:nodedefs.available Operator:LessThan Value:{Quantity:<nil> Percentage:0.1} GracePeriod:0s
MinReclaim:<nil>}}} QOSReserved:map[] ExperimentalCPUManagerPolicy:none ExperimentalCPUManagerReconcilePeriod:10s
ExperimentalPodPidsLimit:-1 EnforceCPULimits:true CPUCFSQuotaPeriod:100ms}
I0330 07:46:25.873200    7713 container_manager_linux.go:286] Creating device plugin manager: true
I0330 07:46:25.873271    7713 state_mem.go:36] [cpumanager] initializing new in-memory state store
I0330 07:46:25.875736    7713 kubelet.go:304] Watching apiserver
I0330 07:46:25.878558    7713 client.go:75] Connecting to docker on unix:///var/run/docker.sock
I0330 07:46:25.878581    7713 client.go:104] Start docker client with request timeout=2m0s
W0330 07:46:25.880560    7713 docker_service.go:561] Hairpin mode set to "promiscuous-bridge" but kubenet is not
enabled, falling back to "hairpin-veth"
I0330 07:46:25.880584    7713 docker_service.go:238] Hairpin mode set to "hairpin-veth"
W0330 07:46:25.880739    7713 cni.go:213] Unable to update cni config: No networks found in /etc/cni/net.d
W0330 07:46:25.882214    7713 hostport_manager.go:68] The binary conntrack is not installed, this can cause
failures in network connection cleanup.
I0330 07:46:25.883177    7713 docker_service.go:253] Docker cri networking managed by kubernetes.io/no-op
```

```

I0330 07:46:25.899118    7713 docker_service.go:258] Docker Info: &
{ID:TQ5X:T6PX:K2WX:LMX6:W44V:RPMB:DME6:AQJP:AMOD:45JW:HCD2:RCXA Containers:0 ContainersRunning:0
ContainersPaused:0 ContainersStopped:0 Images:5 Driver:overlay2 DriverStatus:[[Backing Filesystem extfs] [Supports
d_type true] [Native Overlay Diff true]] SystemStatus:[] Plugins:{Volume:[local] Network:[bridge host macvlan null
overlay] Authorization:[] Log:[awslogs fluentd gcplogs gelf journald json-file local logentries splunk syslog]}
MemoryLimit:true SwapLimit:false KernelMemory:true CPUCfsPeriod:true CPUCfsQuota:true CPUShares:true CPUSet:true
IPv4Forwarding:true BridgeNfIptables:true BridgeNfIP6tables:true Debug:false NFd:22 OomKillDisable:true
NGoroutines:38 SystemTime:2019-03-30T07:46:25.883814796Z LoggingDriver:json-file CgroupDriver:cgroupfs
NEventsListener:0 KernelVersion:4.4.0-1075-aws OperatingSystem:Ubuntu 16.04.5 LTS OSType:linux Architecture:x86_64
IndexServerAddress:https://index.docker.io/v1/ RegistryConfig:0xc00074ae00 NCPU:2 MemTotal:8369913856
GenericResources:[] DockerRootDir:/var/lib/docker HTTPProxy: HTTPSProxy: NoProxy: Name:nodea Labels:[]
ExperimentalBuild:false ServerVersion:18.09.3 ClusterStore: ClusterAdvertise: Runtimes:map[runc:{Path:runc Args:
[]}] DefaultRuntime:runc Swarm:{NodeID: NodeAddr: LocalNodeState:inactive ControlAvailable:false Error:
RemoteManagers:[] Nodes:0 Managers:0 Cluster:<nil>} LiveRestoreEnabled:false Isolation: InitBinary:docker-init
ContainerdCommit:{ID:e6b3f5632f50dbc4e9cb6288d911bf4f5e95b18e Expected:e6b3f5632f50dbc4e9cb6288d911bf4f5e95b18e}
RuncCommit:{ID:6635b4f0c6af3810594d2770f662f34ddc15b40d Expected:6635b4f0c6af3810594d2770f662f34ddc15b40d}
InitCommit:{ID:fec3683 Expected:fec3683} SecurityOptions:[name=apparmor name=seccomp,profile=default]}
I0330 07:46:25.899195    7713 docker_service.go:271] Setting cgroupDriver to cgroupfs
I0330 07:46:25.917301    7713 remote_runtime.go:62] parsed scheme: ""
I0330 07:46:25.917321    7713 remote_runtime.go:62] scheme "" not registered, fallback to default scheme
I0330 07:46:25.917350    7713 remote_image.go:50] parsed scheme: ""
I0330 07:46:25.917356    7713 remote_image.go:50] scheme "" not registered, fallback to default scheme
I0330 07:46:25.917545    7713 asm_amd64.s:1337] ccResolverWrapper: sending new addresses to cc:
[{/var/run/dockershim.sock 0 <nil>}]
I0330 07:46:25.917546    7713 asm_amd64.s:1337] ccResolverWrapper: sending new addresses to cc:
[{/var/run/dockershim.sock 0 <nil>}]
I0330 07:46:25.917559    7713 clientconn.go:796] ClientConn switching balancer to "pick_first"
I0330 07:46:25.917568    7713 clientconn.go:796] ClientConn switching balancer to "pick_first"
I0330 07:46:25.917594    7713 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc0001fda10, CONNECTING
I0330 07:46:25.917599    7713 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc00023bc50, CONNECTING
I0330 07:46:25.918704    7713 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc00023bc50, READY
I0330 07:46:25.919464    7713 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc0001fda10, READY
I0330 07:46:25.922629    7713 kuberuntime_manager.go:210] Container runtime docker initialized, version: 18.09.3,
apiVersion: 1.39.0
I0330 07:46:25.923918    7713 server.go:1037] Started kubelet
E0330 07:46:25.924215    7713 kubelet.go:1282] Image garbage collection failed once. Stats initialization may not
have completed yet: failed to get imageFs info: unable to find data in memory cache
I0330 07:46:25.924716    7713 fs_resource_analyzer.go:64] Starting FS ResourceAnalyzer
I0330 07:46:25.924746    7713 status_manager.go:152] Starting to sync pod status with apiserver

```

```

I0330 07:46:25.924761    7713 kubelet.go:1806] Starting kubelet main sync loop.
I0330 07:46:25.924778    7713 kubelet.go:1823] skipping pod synchronization - [container runtime status check may
not have completed yet., PLEG is not healthy: pleg has yet to be successful.]
I0330 07:46:25.924865    7713 server.go:141] Starting to listen on 0.0.0.0:10250
I0330 07:46:25.925483    7713 server.go:343] Adding debug handlers to kubelet server.
I0330 07:46:25.927000    7713 volume_manager.go:248] Starting Kubelet Volume Manager
I0330 07:46:25.928060    7713 desired_state_of_world_populator.go:130] Desired state populator starts to run
E0330 07:46:25.948811    7713 controller.go:194] failed to get node "nodea" when trying to set owner ref to the
node lease: nodes "nodea" not found
I0330 07:46:25.953124    7713 clientconn.go:440] parsed scheme: "unix"
I0330 07:46:25.953277    7713 clientconn.go:440] scheme "unix" not registered, fallback to default scheme
I0330 07:46:25.953312    7713 asm_amd64.s:1337] ccResolverWrapper: sending new addresses to cc:
[{"unix:///run/containerd/containerd.sock 0 <nil>}]
I0330 07:46:25.953327    7713 clientconn.go:796] ClientConn switching balancer to "pick_first"
I0330 07:46:25.953363    7713 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc000b14660, CONNECTING
I0330 07:46:25.953517    7713 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc000b14660, READY
I0330 07:46:26.023125    7713 kubelet_node_status.go:283] Setting node annotation to enable volume controller
attach/detach
I0330 07:46:26.025001    7713 cpu_manager.go:155] [cpumanager] starting with none policy
I0330 07:46:26.025018    7713 cpu_manager.go:156] [cpumanager] reconciling every 10s
I0330 07:46:26.025034    7713 policy_none.go:42] [cpumanager] none policy: Start
W0330 07:46:26.025633    7713 manager.go:538] Failed to retrieve checkpoint for "kubelet_internal_checkpoint":
checkpoint is not found
W0330 07:46:26.026074    7713 container_manager_linux.go:818] CPUAccounting not enabled for pid: 7713
W0330 07:46:26.026090    7713 container_manager_linux.go:821] MemoryAccounting not enabled for pid: 7713
E0330 07:46:26.026520    7713 eviction_manager.go:247] eviction manager: failed to get summary stats: failed to
get node info: node "nodea" not found
I0330 07:46:26.038838    7713 kubelet_node_status.go:283] Setting node annotation to enable volume controller
attach/detach
W0330 07:46:26.039128    7713 pod_container_deletor.go:75] Container
"85176a97712ee4859924d45d113c4e25a8c8828a3a456409841df57633f8b8c4" not found in pod's containers
E0330 07:46:26.039165    7713 kubelet.go:2244] node "nodea" not found
I0330 07:46:26.040942    7713 kubelet_node_status.go:72] Attempting to register node nodea
I0330 07:46:26.043668    7713 kubelet_node_status.go:75] Successfully registered node nodea
I0330 07:46:26.138857    7713 reconciler.go:154] Reconciler: start to sync state

```

...

Verify the cluster (with one node so far). Before we can use the `kubectl` command we need to specify the cluster we want to interact with, again substitute your cluster master IP in the example below:

```
ubuntu@nodea:~$ kubectl config set-cluster local --server=http://nodea:8080

Cluster "local" set.
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl config set-context local --cluster=local

Context "local" created.
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl config use-context local

Switched to context "local".
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl get nodes

NAME      STATUS    ROLES    AGE   VERSION
nodea     Ready     <none>   81s   v1.14.0
ubuntu@nodea:~$
```

Now let recreate our simple Pod on *nodea* (from lab 1), as a reminder, this is the yaml:

```
ubuntu@nodea:~$ cat testpod.yaml

apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  nodeName: nodea
  automountServiceAccountToken: false
  containers:
  - name: nginx
```

```
image: nginx
ports:
- containerPort: 80
volumeMounts:
- mountPath: /var/log/nginx
  name: nginx-logs
- name: log-truncator
  image: busybox
  command:
  - /bin/sh
  args: [-c, 'while true; do cat /dev/null > /logdir/access.log; sleep 10; done']
  volumeMounts:
  - mountPath: /logdir
    name: nginx-logs
volumes:
- name: nginx-logs
  emptyDir: {}
ubuntu@nodea:~$
```

Deploy your pod via create subcommand.

```
ubuntu@nodea:~$ kubectl create -f testpod.yaml

pod/nginx created
ubuntu@nodea:~$
```

Confirm your pod has entered the *Running* state via `kubectl get pod`.

```
ubuntu@nodea:~$ kubectl get pods

NAME      READY   STATUS    RESTARTS   AGE
nginx     2/2     Running   0           18s
ubuntu@nodea:~$
```

We will now locate the node our pod has been deployed to (remember, we have not added *nodeb* to the cluster, yet).

```
ubuntu@nodea:~$ kubectl describe pod nginx | grep -E ^Node:
```

```
Node:          nodea/172.31.28.198
ubuntu@nodea:~$
```

or

```
ubuntu@nodea:~$ curl -s http://localhost:8080/api/v1/pods | jq .items[].spec.nodeName -r
nodea
ubuntu@nodea:~$
```

If you review our pods template, you will notice an entry *spec.nodeName*. This field is where we hardcoded the node where our pod was placed.

## 7. Run a pod without *nodeName*

Open copy `testpod.yaml`, change the pod name to `nginx-a` and remove the option *nodeName*; leave everything else the same as before.

```
ubuntu@nodea:~$ cp testpod.yaml testpod-a.yaml
ubuntu@nodea:~$ vim testpod-a.yaml
ubuntu@nodea:~$ cat testpod-a.yaml

apiVersion: v1
kind: Pod
metadata:
  name: nginx-a
spec:
#  nodeName: nodea
  containers:

...

ubuntu@nodea:~$
```

Launch the pod again and monitor its status.

```
ubuntu@nodea:~$ kubectl create -f testpod-a.yaml

pod/nginx-a created
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl get pods

NAME      READY   STATUS    RESTARTS   AGE
nginx     2/2     Running   0           3m
nginx-a   0/2     Pending   0           9s
ubuntu@nodea:~$
```

Notice, the status is "Pending". Why?

The pod has no target host, which means that it must be scheduled to a node but we have no scheduler!

The pod will remain in the pending state until you either recreate the pod with a *nodeName* configured, or start the scheduler.

## 8. Add nodeb to the cluster

Before we start the scheduler let's add nodeb to the cluster. To do this, we need to install the `kubelet` services on *nodeb*. Since we have already compiled it on *nodea* we will simply copy it (and everything else) over.

**On nodeb** run the following commands:

```
laptop$ ssh -i k8s-adv-student.pem ubuntu@<external-ip>

...

ubuntu@nodeb:~$ mkdir kube-bin

ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ nc -l -p 7000 | tar xv -C kube-bin/
```

...

This puts netcat in listening mode, with tar decompressing into the `~/kube-bin` directory.

**On nodea** run (it will take a couple minutes to complete):

First add nodeb to the etc/hosts file:

```
ubuntu@nodea:~$ sudo vim /etc/hosts

ubuntu@nodea:~$ cat /etc/hosts

127.0.0.1 localhost
172.31.30.148 nodeb
172.31.28.198 nodea

# The following lines are desirable for IPv6 capable hosts
::1 ip6-localhost ip6-loopback
fe00::0 ip6-localnet
ff00::0 ip6-mcastprefix
ff02::1 ip6-allnodes
ff02::2 ip6-allrouters
ff02::3 ip6-allhosts
ubuntu@nodea:~$
```

Now tar the Kubernetes binaries into the netcat tunnel (it will take a couple minutes to complete):

```
ubuntu@nodea:~$ tar -C ~/k8s/_output/local/bin/linux/amd64/ -cf - . | nc nodeb 7000

ubuntu@nodea:~$
```

This command uses netcat to funnel our tared data over to nodeb. This will copy our binaries from *nodea* to *nodeb* (you will see output on *nodeb*). Depending on what you compiled, your output may differ slightly on nodeb.

```
ubuntu@nodeb:~$ nc -l -p 7000 | tar xv -C kube-bin/
```



```
./
./gendocs
./genman
./kube-apiserver
./genswaggertypedocs
./linkcheck
./conversion-gen
./teststale
./go-bindata
./defaulter-gen
./genyaml
./hyperkube
./kube-aggregator
./deepcopy-gen
./genfeddocs
./kubelet
./kube-proxy
./genkubedocs
./kubeadm
./kube-scheduler
./gke-certificates-controller
./kube-controller-manager
./mungedocs
./apiextensions-apiserver
./openapi-gen
./cloud-controller-manager
./ginkgo
./e2e.test
./kubemark
./kubectl
./e2e_node.test
./kubefed
```

```
ubuntu@nodeb:~$
```

Before running the kubelet on nodeb, create a kubeconfig file with information to connect to the nodea apiserver.

```
ubuntu@nodeb:~$ vim nodeb.conf
ubuntu@nodeb:~$ cat nodeb.conf
```

```
apiVersion: v1
clusters:
```

```
- cluster:
  server: http://nodea:8080
  name: local
contexts:
- context:
  cluster: local
  user: ""
  name: local
current-context: local
kind: Config
preferences: {}
users: []
ubuntu@nodeb:~$
```

Note that all our kubelet really needs to know is the URI of the API server.

```
ubuntu@nodeb:~$ vim nodeb.yaml
ubuntu@nodeb:~$ cat nodeb.yaml

apiVersion: kubelet.config.k8s.io/v1beta1
kind: KubeletConfiguration
authentication:
  anonymous:
    enabled: true
cgroupDriver: cgroupfs
failSwapOn: true
ubuntu@nodeb:~$
```

On *nodeb*, you can start the `kubelet` process via:

```
ubuntu@nodeb:~$ sudo $HOME/kube-bin/kubelet \
--kubeconfig=nodeb.conf \
--config=nodeb.yaml \
--allow-privileged=true \
--runtime-cgroups=/systemd/machine.slice \
--kubelet-cgroups=/systemd/machine.slice \
--pod-infra-container-image=k8s.gcr.io/pause:3.1
```

Flag `--allow-privileged` has been deprecated, will be removed in a future version

Flag `--kubelet-cgroups` has been deprecated, This parameter should be set via the config file specified by the

Kubelet's --config flag. See <https://kubernetes.io/docs/tasks/administer-cluster/kubelet-config-file/> for more information.

Flag --allow-privileged has been deprecated, will be removed in a future version

Flag --kubelet-cgroups has been deprecated, This parameter should be set via the config file specified by the Kubelet's --config flag. See <https://kubernetes.io/docs/tasks/administer-cluster/kubelet-config-file/> for more information.

```

I0330 08:19:31.687354    4673 server.go:417] Version: v1.14.0
I0330 08:19:31.687718    4673 plugins.go:103] No cloud provider specified.
I0330 08:19:31.731090    4673 server.go:625] --cgroups-per-qos enabled, but --cgroup-root was not specified.
defaulting to /
I0330 08:19:31.731380    4673 container_manager_linux.go:261] container manager verified user specified cgroup-
root exists: []
I0330 08:19:31.731401    4673 container_manager_linux.go:266] Creating Container Manager object based on Node
Config: {RuntimeCgroupsName:/systemd/machine.slice SystemCgroupsName: KubeletCgroupsName:/systemd/machine.slice
ContainerRuntime:docker CgroupsPerQOS:true CgroupRoot:/ CgroupDriver:cgroupfs KubeletRootDir:/var/lib/kubelet
ProtectKernelDefaults:false NodeAllocatableConfig:{KubeReservedCgroupName: SystemReservedCgroupName:
EnforceNodeAllocatable:map[pods:{}] KubeReserved:map[] SystemReserved:map[] HardEvictionThresholds:
[{Signal:memory.available Operator:LessThan Value:{Quantity:100Mi Percentage:0} GracePeriod:0s MinReclaim:<nil>}
{Signal:nodefs.available Operator:LessThan Value:{Quantity:<nil> Percentage:0.1} GracePeriod:0s MinReclaim:<nil>}
{Signal:nodefs.inodesFree Operator:LessThan Value:{Quantity:<nil> Percentage:0.05} GracePeriod:0s MinReclaim:
<nil>} {Signal:imagefs.available Operator:LessThan Value:{Quantity:<nil> Percentage:0.15} GracePeriod:0s
MinReclaim:<nil>}}} QOSReserved:map[] ExperimentalCPUManagerPolicy:none ExperimentalCPUManagerReconcilePeriod:10s
ExperimentalPodPidsLimit:-1 EnforceCPULimits:true CPUCFSQuotaPeriod:100ms}
I0330 08:19:31.731499    4673 container_manager_linux.go:286] Creating device plugin manager: true
I0330 08:19:31.731570    4673 state_mem.go:36] [cpumanager] initializing new in-memory state store
I0330 08:19:31.738047    4673 kubelet.go:304] Watching apiserver
I0330 08:19:31.740203    4673 client.go:75] Connecting to docker on unix:///var/run/docker.sock
I0330 08:19:31.740227    4673 client.go:104] Start docker client with request timeout=2m0s
W0330 08:19:31.741546    4673 docker_service.go:561] Hairpin mode set to "promiscuous-bridge" but kubenet is not
enabled, falling back to "hairpin-veth"
I0330 08:19:31.741569    4673 docker_service.go:238] Hairpin mode set to "hairpin-veth"
W0330 08:19:31.741662    4673 cni.go:213] Unable to update cni config: No networks found in /etc/cni/net.d
W0330 08:19:31.743077    4673 hostport_manager.go:68] The binary conntrack is not installed, this can cause
failures in network connection cleanup.
I0330 08:19:31.744096    4673 docker_service.go:253] Docker cri networking managed by kubernetes.io/no-op
I0330 08:19:31.761066    4673 docker_service.go:258] Docker Info: &
{ID:6TFI:7NRU:W6A5:B2NU:VOXK:OYQG:FJNE:4MSB:KG5T:P70Q:SMJW:ATUN Containers:0 ContainersRunning:0
ContainersPaused:0 ContainersStopped:0 Images:0 Driver:overlay2 DriverStatus:[{Backing Filesystem extfs} {Supports
d_type true} {Native Overlay Diff true}] SystemStatus:[] Plugins:{Volume:[local] Network:[bridge host macvlan null
overlay] Authorization:[] Log:[awslogs fluentd gcplogs gelf journald json-file local logentries splunk syslog]}
MemoryLimit:true SwapLimit:false KernelMemory:true CPUCfsPeriod:true CPUCfsQuota:true CPUShares:true CPUSet:true
IPv4Forwarding:true BridgeNfIptables:true BridgeNfIP6tables:true Debug:false NFd:22 OomKillDisable:true
NGoroutines:37 SystemTime:2019-03-30T08:19:31.744762922Z LoggingDriver:json-file CgroupDriver:cgroupfs

```

```

NEventsListener:0 KernelVersion:4.4.0-1075-aws OperatingSystem:Ubuntu 16.04.5 LTS OSType:linux Architecture:x86_64
IndexServerAddress:https://index.docker.io/v1/ RegistryConfig:0xc0007838f0 NCPU:2 MemTotal:4142067712
GenericResources:[] DockerRootDir:/var/lib/docker HTTPProxy: HTTPSProxy: NoProxy: Name:nodeb Labels:[]
ExperimentalBuild:false ServerVersion:18.09.4 ClusterStore: ClusterAdvertise: Runtimes:map[runc:{Path:runc Args:
[]}] DefaultRuntime:runc Swarm:{NodeID: NodeAddr: LocalNodeState:inactive ControlAvailable:false Error:
RemoteManagers:[] Nodes:0 Managers:0 Cluster:<nil>} LiveRestoreEnabled:false Isolation: InitBinary:docker-init
ContainerdCommit:{ID:bb71b10fd8f58240ca47fbb579b9d1028eea7c84 Expected:bb71b10fd8f58240ca47fbb579b9d1028eea7c84}
RuncCommit:{ID:2b18fe1d885ee5083ef9f0838fee39b62d653e30 Expected:2b18fe1d885ee5083ef9f0838fee39b62d653e30}
InitCommit:{ID:fec3683 Expected:fec3683} SecurityOptions:[name=apparmor name=seccomp,profile=default]}
I0330 08:19:31.761152 4673 docker_service.go:271] Setting cgroupDriver to cgroupfs
I0330 08:19:31.778932 4673 remote_runtime.go:62] parsed scheme: ""
I0330 08:19:31.778956 4673 remote_runtime.go:62] scheme "" not registered, fallback to default scheme
I0330 08:19:31.778986 4673 remote_image.go:50] parsed scheme: ""
I0330 08:19:31.778995 4673 remote_image.go:50] scheme "" not registered, fallback to default scheme
I0330 08:19:31.779042 4673 asm_amd64.s:1337] ccResolverWrapper: sending new addresses to cc:
[{/var/run/dockershim.sock 0 <nil>}]
I0330 08:19:31.779058 4673 clientconn.go:796] ClientConn switching balancer to "pick_first"
I0330 08:19:31.779098 4673 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc0001d3580, CONNECTING
I0330 08:19:31.779133 4673 asm_amd64.s:1337] ccResolverWrapper: sending new addresses to cc:
[{/var/run/dockershim.sock 0 <nil>}]
I0330 08:19:31.779143 4673 clientconn.go:796] ClientConn switching balancer to "pick_first"
I0330 08:19:31.779170 4673 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc0001d8680, CONNECTING
I0330 08:19:31.779694 4673 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc0001d3580, READY
I0330 08:19:31.781156 4673 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc0001d8680, READY
I0330 08:19:31.784413 4673 kuberuntime_manager.go:210] Container runtime docker initialized, version: 18.09.4,
apiVersion: 1.39.0
W0330 08:19:31.784610 4673 probe.go:268] Flexvolume plugin directory at /usr/libexec/kubernetes/kubelet-
plugins/volume/exec/ does not exist. Recreating.
I0330 08:19:31.786082 4673 server.go:1037] Started kubelet
E0330 08:19:31.786529 4673 kubelet.go:1282] Image garbage collection failed once. Stats initialization may not
have completed yet: failed to get imageFs info: unable to find data in memory cache
I0330 08:19:31.787175 4673 fs_resource_analyzer.go:64] Starting FS ResourceAnalyzer
I0330 08:19:31.787264 4673 status_manager.go:152] Starting to sync pod status with apiserver
I0330 08:19:31.787356 4673 kubelet.go:1806] Starting kubelet main sync loop.
I0330 08:19:31.787436 4673 kubelet.go:1823] skipping pod synchronization - [container runtime status check may
not have completed yet., PLEG is not healthy: pleg has yet to be successful.]
I0330 08:19:31.787581 4673 server.go:141] Starting to listen on 0.0.0.0:10250
I0330 08:19:31.788282 4673 server.go:343] Adding debug handlers to kubelet server.
I0330 08:19:31.789427 4673 volume_manager.go:248] Starting Kubelet Volume Manager

```

```

E0330 08:19:31.798544    4673 controller.go:194] failed to get node "nodeb" when trying to set owner ref to the
node lease: nodes "nodeb" not found
I0330 08:19:31.800275    4673 desired_state_of_world_populator.go:130] Desired state populator starts to run
I0330 08:19:31.824668    4673 clientconn.go:440] parsed scheme: "unix"
I0330 08:19:31.824822    4673 clientconn.go:440] scheme "unix" not registered, fallback to default scheme
I0330 08:19:31.824949    4673 asm_amd64.s:1337] ccResolverWrapper: sending new addresses to cc:
[unix:///run/containerd/containerd.sock 0 <nil>}]
I0330 08:19:31.824967    4673 clientconn.go:796] ClientConn switching balancer to "pick_first"
I0330 08:19:31.825002    4673 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc000235bf0, CONNECTING
I0330 08:19:31.825270    4673 balancer_conn_wrappers.go:131] pickfirstBalancer: HandleSubConnStateChange:
0xc000235bf0, READY
E0330 08:19:31.895713    4673 kubelet.go:2244] node "nodeb" not found
I0330 08:19:31.895740    4673 kubelet.go:1823] skipping pod synchronization - container runtime status check may
not have completed yet.
I0330 08:19:31.895768    4673 kubelet_node_status.go:283] Setting node annotation to enable volume controller
attach/detach
I0330 08:19:31.897665    4673 kubelet_node_status.go:72] Attempting to register node nodeb
I0330 08:19:31.901600    4673 kubelet_node_status.go:75] Successfully registered node nodeb
I0330 08:19:31.905029    4673 cpu_manager.go:155] [cpumanager] starting with none policy
I0330 08:19:31.905047    4673 cpu_manager.go:156] [cpumanager] reconciling every 10s
I0330 08:19:31.905065    4673 policy_none.go:42] [cpumanager] none policy: Start
W0330 08:19:31.917027    4673 manager.go:538] Failed to retrieve checkpoint for "kubelet_internal_checkpoint":
checkpoint is not found
W0330 08:19:31.920491    4673 container_manager_linux.go:818] CPUAccounting not enabled for pid: 4673
W0330 08:19:31.920509    4673 container_manager_linux.go:821] MemoryAccounting not enabled for pid: 4673
I0330 08:19:32.100782    4673 reconciler.go:154] Reconciler: start to sync state

```

...

To confirm the *nodeb* `kubelet` has connected to the kube-apiserver on *nodeb*, run the following commands on *nodeb* to configure the kubelet and get the cluster node list.

```
ubuntu@nodeb:~$ sudo cp ./kube-bin/kubect1 /usr/bin/
```

```
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ kubect1 config set-cluster local --server=http://nodea:8080
```

```
Cluster "local" set.  
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ kubectl config set-context local --cluster=local  
  
Context "local" created.  
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ kubectl config use-context local  
  
Switched to context "local".  
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ kubectl get nodes  
  
NAME      STATUS    ROLES    AGE   VERSION  
nodea     Ready     <none>   35m   v1.14.0  
nodeb     Ready     <none>   2m4s  v1.14.0  
ubuntu@nodeb:~$
```

You can also use `curl` (with help from `jq`) directly against the API:

```
ubuntu@nodeb:~$ sudo apt-get -y install jq  
  
...  
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ curl -s http://nodea:8080/api/v1/nodes | jq -r .items[].metadata.name  
  
nodea  
nodeb  
ubuntu@nodeb:~$
```

or to see the full output:

```
ubuntu@nodeb:~$ curl -s http://nodea:8080/api/v1/nodes
{
  "kind": "NodeList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/nodes",
    "resourceVersion": "264"
  },
  "items": [
    {
      "metadata": {
        "name": "nodea",
        ...
      }
    }
  ]
}
```

## 9. Running a pod on *nodeb*

Back on *nodea*, copy and modify the `testpod.yaml` to include:

```
ubuntu@nodea:~$ cp testpod-a.yaml testpod-b.yaml

ubuntu@nodea:~$ vim testpod-b.yaml
ubuntu@nodea:~$ cat testpod-b.yaml

apiVersion: v1
kind: Pod
metadata:
  name: nginx-b
spec:
  nodeName: nodeb
  containers:
  ...
```

Launch the pod, as you proceed check status with the following methods.

- via `kubectl`
- via `docker`
- via `curl`

If you see status "ContainerCreating", this typically indicates the node is pulling the container image. Recall that *nodeb* is a brand new Docker install and as of yet has no local images to work with. Each node must pull its own images.

```
ubuntu@nodea:~$ kubectl create -f testpod-b.yaml

pod/nginx-b created
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	2/2	Running	0	18m
nginx-a	0/2	Pending	0	15m
nginx-b	0/2	ContainerCreating	0	7s

```
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	2/2	Running	0	5m
nginx-a	0/2	Pending	0	56s
nginx-b	2/2	Running	0	25s

```
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl get pod nginx-b \
-o=custom-columns=Name:.metadata.name,hostIP:.status.hostIP
```



```
Name      hostIP
nginx-b    172.31.30.148
ubuntu@nodea:~$
```

On nodeb:

```
ubuntu@nodeb:~$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
885652b2eeb1	busybox	"/bin/sh -c 'while t..."	About a minute ago	Up About a minute
k8s_log-truncator_nginx_default_bea2bed7-4cef-11e8-8645-000c29473113_0	nginx	"nginx -g 'daemon of..."	2 minutes ago	Up 2 minutes
k8s_nginx_nginx_default_bea2bed7-4cef-11e8-8645-000c29473113_0	k8s.gcr.io/pause-amd64:3.1	"/pause"	3 minutes ago	Up 3 minutes
b319ff91f9c4				
k8s_POD_nginx_default_bea2bed7-4cef-11e8-8645-000c29473113_0				

```
ubuntu@nodeb:~$
```

```
ubuntu@nodeb:~$ curl -s http://nodea:8080/api/v1/pods
```

```
{
  "kind": "PodList",
  "apiVersion": "v1",
  "metadata": {
    "selfLink": "/api/v1/pods",
    "resourceVersion": "213"
  },
  "items": [
    {
      "metadata": {
        "name": "nginx",
        "namespace": "default",
        "selfLink": "/api/v1/namespaces/default/pods/nginx",
        "uid": "4c12afc0-fbd8-11e7-bf93-000c290928aa",
        "resourceVersion": "194",
        "creationTimestamp": "2018-01-17T22:46:46Z"
      },

```

```
...
```

```
ubuntu@nodeb:~$
```

## 10. Start the scheduler

To run the scheduler we can simply execute the binary with a switch pointing it to the api server. **In a new terminal** on nodea:

```
laptop$ ssh -i k8s-adv-student.pem ubuntu@<external-ip>
```

```
...
```

```
ubuntu@nodea:~$ $HOME/k8s/_output/bin/kube-scheduler --kubeconfig=nodea.conf
```

```
I0330 08:31:14.877734    9479 serving.go:319] Generated self-signed cert in-memory
W0330 08:31:15.494975    9479 authentication.go:249] No authentication-kubeconfig provided in order to lookup
client-ca-file in configmap/extension-apiserver-authentication in kube-system, so client certificate
authentication won't work.
W0330 08:31:15.495002    9479 authentication.go:252] No authentication-kubeconfig provided in order to lookup
requestheader-client-ca-file in configmap/extension-apiserver-authentication in kube-system, so request-header
client certificate authentication won't work.
W0330 08:31:15.495015    9479 authorization.go:146] No authorization-kubeconfig provided, so SubjectAccessReview
of authorization tokens won't work.
I0330 08:31:15.497506    9479 server.go:142] Version: v1.14.0
I0330 08:31:15.497574    9479 defaults.go:87] TaintNodesByCondition is enabled, PodToleratesNodeTaints predicate
is mandatory
W0330 08:31:15.499062    9479 authorization.go:47] Authorization is disabled
W0330 08:31:15.499078    9479 authentication.go:55] Authentication is disabled
I0330 08:31:15.499092    9479 deprecated_insecure_serving.go:49] Serving healthz insecurely on [::]:10251
I0330 08:31:15.499503    9479 secure_serving.go:116] Serving securely on [::]:10259
I0330 08:31:16.401540    9479 controller_utils.go:1027] Waiting for caches to sync for scheduler controller
I0330 08:31:16.501749    9479 controller_utils.go:1034] Caches are synced for scheduler controller
I0330 08:31:16.501828    9479 leaderelection.go:217] attempting to acquire leader lease kube-system/kube-
scheduler...
I0330 08:31:16.506049    9479 leaderelection.go:227] successfully acquired lease kube-system/kube-scheduler
```

Note that while many API servers can run in parallel (etcd ensures state is always consistent) only one scheduler may run within the cluster to avoid scheduling

conflicts. For this reason the scheduler causes an election using etcd to determine which of the possibly several schedulers running will become the leader. All other schedulers simply monitor the leader for failure. If the leader fails, the remaining schedulers elect an new leader.

## 11. Pod placement via the scheduler

Now that we have built and started the scheduler, check in with our "pending" pods:

```
ubuntu@nodea:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	2/2	Running	0	52m
nginx-a	0/2	Pending	0	45m
nginx-b	2/2	Running	0	17m

```
ubuntu@nodea:~$
```

What happened?

In Kubernetes terms, the nodes are tainted. A taint consists of a *key*, a *value*, and an *effect*. The effect must be *NoSchedule*, *PreferNoSchedule* or *NoExecute*. You can view the taints on your node with the `kubectl` command. Use the `kubectl describe` subcommand to see details for one of your nodes:

```
ubuntu@nodea:~$ kubectl describe node nodea | grep Taints
```

Taints: node.kubernetes.io/not-ready:NoSchedule

```
ubuntu@nodea:~$
```

This means the `kube-scheduler` can not place pods on this node. To remove this taint we can use the `kubectl taint` subcommand.

**NOTE** The command below removes ("-") the taint from all (--all) nodes in the cluster. **Do not forget the trailing -** The `-` is what tells Kubernetes to remove the taint!

We know what you're thinking and we agree, "taint" is an awful name for this feature and a trailing dash with no space is an equally wacky way to remove something.

```
ubuntu@nodea:~$ kubectl taint nodes --all node.kubernetes.io/not-ready-
```

```
node/nodea untainted
node/nodeb untainted
```

```
ubuntu@nodea:~$ kubectl describe node nodea | grep Taints
```

```
Taints:                <none>
ubuntu@nodea:~$
```

Check in with our "pending" pods once more:

```
ubuntu@nodea:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	2/2	Running	0	70m
nginx-a	2/2	Running	0	64m
nginx-b	2/2	Running	0	36m

```
ubuntu@nodea:~$
```

Where does the pod land?

```
ubuntu@nodea:~$ kubectl get pod nginx-a \
-o=custom-columns=Name:.metadata.name,hostIP:.status.hostIP
```

Name	hostIP
nginx-a	172.31.28.198

```
ubuntu@nodea:~$
```

## 12. Launch additional pods

Create another pod with a random name.

```
ubuntu@nodea:~$ sed -e '/nodeName/d' testpod-a.yaml \
-e "s/name: nginx/name: nginx-$RANDOM/g" | kubectl create -f -
```

```
pod/nginx-29538-a created
ubuntu@nodea:~$
```

Determine which Node is our new pod running on:

```
ubuntu@nodea:~$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
nginx	2/2	Running	0	71m	172.17.0.2	nodea	<none>		<none>	
nginx-12095-a	2/2	Running	0	10s	172.17.0.3	nodeb	<none>		<none>	
nginx-a	2/2	Running	0	65m	172.17.0.3	nodea	<none>		<none>	
nginx-b	2/2	Running	0	36m	172.17.0.2	nodeb	<none>		<none>	

```
ubuntu@nodea:~$
```

Create several more pods and view which node a pod is placed on. The default scheduler will spread pods across the nodes.

```
ubuntu@nodea:~$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
nginx	2/2	Running	0	72m	172.17.0.2	nodea	<none>		<none>	
nginx-11271-a	2/2	Running	0	17s	172.17.0.4	nodea	<none>		<none>	
nginx-11427-a	2/2	Running	0	16s	172.17.0.5	nodeb	<none>		<none>	
nginx-12095-a	2/2	Running	0	58s	172.17.0.3	nodeb	<none>		<none>	
nginx-12751-a	2/2	Running	0	19s	172.17.0.4	nodeb	<none>		<none>	
nginx-a	2/2	Running	0	65m	172.17.0.3	nodea	<none>		<none>	
nginx-b	2/2	Running	0	37m	172.17.0.2	nodeb	<none>		<none>	

```
ubuntu@nodea:~$
```

Try to remove all the pods (**note: your terminal should hang at this command**):

```
ubuntu@nodea:~$ kubectl get pod -o go-template \
--template '{{range .items}}{{.metadata.name}}{{"\n"}}{{end}}' | xargs kubectl delete pod

pod "nginx" deleted
pod "nginx-23152-a" deleted
pod "nginx-25624-a" deleted
```

```
pod "nginx-25979-a" deleted
pod "nginx-29538-a" deleted
pod "nginx-9358-a" deleted
pod "nginx-a" deleted
pod "nginx-b" deleted
```

What's going on here?

Some of the Pod clean up tasks are performed by the controller manager, notice how the pods hang on deletion. By running the controller manager we can resolve the pod shutdown.

## 13. Starting the controller manager

The old approach to running non-API Server Kubernetes services was to provide them with the IRI of the API Server on the command line. The kube config approach is the go forward way to centralize cluster configuration for all of the Kubernetes services on a given node. Command line switches can still be used for many features (and are still required in some cases) but for basic operation the kubeconfig should suffice.

Run the controller manager and use the kubeconfig created earlier to point the controller manager at the appropriate API server; as a reminder, this is what the kubeconfig looks like:

```
ubuntu@nodea:~$ cat nodea.conf

apiVersion: v1
clusters:
- cluster:
    server: http://nodea:8080
    name: local
contexts:
- context:
    cluster: local
    user: ""
    name: local
current-context: local
kind: Config
preferences: {}
users: []
ubuntu@nodea:~$
```

**In a new terminal** start the controller manager:

```
ubuntu@nodea:~$ $HOME/k8s/_output/bin/kube-controller-manager --kubeconfig=nodea.conf
```

```
I0330 19:57:07.651811 13668 serving.go:319] Generated self-signed cert in-memory
W0330 19:57:08.117032 13668 authentication.go:249] No authentication-kubeconfig provided in order to lookup
client-ca-file in configmap/extension-apiserver-authentication in kube-system, so client certificate
authentication won't work.
W0330 19:57:08.117061 13668 authentication.go:252] No authentication-kubeconfig provided in order to lookup
requestheader-client-ca-file in configmap/extension-apiserver-authentication in kube-system, so request-header
client certificate authentication won't work.
W0330 19:57:08.117075 13668 authorization.go:146] No authorization-kubeconfig provided, so SubjectAccessReview
of authorization tokens won't work.
I0330 19:57:08.117116 13668 controllermanager.go:155] Version: v1.14.0
I0330 19:57:08.117556 13668 secure_serving.go:116] Serving securely on [::]:10257
I0330 19:57:08.117982 13668 deprecated_insecure_serving.go:51] Serving insecurely on [::]:10252
I0330 19:57:08.118091 13668 leaderelection.go:217] attempting to acquire leader lease kube-system/kube-
controller-manager...
I0330 19:57:08.122732 13668 leaderelection.go:227] successfully acquired lease kube-system/kube-controller-
manager
I0330 19:57:08.123112 13668 event.go:209] Event(v1.ObjectReference{Kind:"Endpoints", Namespace:"kube-system",
Name:"kube-controller-manager", UID:"ffc05c46-5325-11e9-92a6-02ef63d53bbe", APIVersion:"v1",
ResourceVersion:"785", FieldPath:""}): type: 'Normal' reason: 'LeaderElection' nodea_ffc005c8-5325-11e9-bf69-
02ef63d53bbe became leader
I0330 19:57:08.334487 13668 plugins.go:103] No cloud provider specified.
W0330 19:57:08.334531 13668 controllermanager.go:517] "serviceaccount-token" is disabled because there is no
private key
I0330 19:57:08.335193 13668 controllermanager.go:497] Started "disruption"
I0330 19:57:08.335411 13668 node_lifecycle_controller.go:292] Sending events to api server.
I0330 19:57:08.335488 13668 disruption.go:286] Starting disruption controller
I0330 19:57:08.335559 13668 node_lifecycle_controller.go:325] Controller is using taint based evictions.
I0330 19:57:08.335566 13668 controller_utils.go:1027] Waiting for caches to sync for disruption controller
I0330 19:57:08.335613 13668 taint_manager.go:175] Sending events to api server.
I0330 19:57:08.335996 13668 node_lifecycle_controller.go:390] Controller will reconcile labels.
I0330 19:57:08.336021 13668 node_lifecycle_controller.go:403] Controller will taint node by condition.
I0330 19:57:08.336083 13668 controllermanager.go:497] Started "nodelifecycle"
I0330 19:57:08.336485 13668 controllermanager.go:497] Started "persistentvolume-expander"
I0330 19:57:08.336745 13668 node_lifecycle_controller.go:427] Starting node controller
I0330 19:57:08.336778 13668 controller_utils.go:1027] Waiting for caches to sync for taint controller
I0330 19:57:08.336811 13668 expand_controller.go:153] Starting expand controller
I0330 19:57:08.336829 13668 controller_utils.go:1027] Waiting for caches to sync for expand controller
I0330 19:57:08.336882 13668 controllermanager.go:497] Started "pvc-protection"
I0330 19:57:08.336890 13668 pvc_protection_controller.go:99] Starting PVC protection controller
I0330 19:57:08.336920 13668 controller_utils.go:1027] Waiting for caches to sync for PVC protection controller
```

```

I0330 19:57:08.337482 13668 controllermanager.go:497] Started "horizontalpodautoscaling"
I0330 19:57:08.337657 13668 horizontal.go:156] Starting HPA controller
I0330 19:57:08.337678 13668 controller_utils.go:1027] Waiting for caches to sync for HPA controller
I0330 19:57:08.337884 13668 controllermanager.go:497] Started "job"
W0330 19:57:08.337902 13668 controllermanager.go:476] "bootstrapsigner" is disabled
I0330 19:57:08.337912 13668 job_controller.go:143] Starting job controller
I0330 19:57:08.337933 13668 controller_utils.go:1027] Waiting for caches to sync for job controller
I0330 19:57:08.361029 13668 controllermanager.go:497] Started "namespace"
I0330 19:57:08.361119 13668 namespace_controller.go:186] Starting namespace controller
I0330 19:57:08.361180 13668 controller_utils.go:1027] Waiting for caches to sync for namespace controller
I0330 19:57:08.361472 13668 controllermanager.go:497] Started "deployment"
W0330 19:57:08.361488 13668 controllermanager.go:476] "tokencleaner" is disabled
I0330 19:57:08.361809 13668 deployment_controller.go:152] Starting deployment controller
I0330 19:57:08.361875 13668 controller_utils.go:1027] Waiting for caches to sync for deployment controller
I0330 19:57:08.361827 13668 controllermanager.go:497] Started "clusterrole-aggregation"
I0330 19:57:08.361834 13668 clusterroleaggregation_controller.go:148] Starting ClusterRoleAggregator
I0330 19:57:08.362238 13668 controller_utils.go:1027] Waiting for caches to sync for ClusterRoleAggregator
controller
I0330 19:57:08.767918 13668 controllermanager.go:497] Started "garbagecollector"
W0330 19:57:08.767967 13668 controllermanager.go:489] Skipping "csrsigning"
I0330 19:57:08.768383 13668 controllermanager.go:497] Started "ttl"
W0330 19:57:08.768402 13668 controllermanager.go:489] Skipping "ttl-after-finished"
I0330 19:57:08.768883 13668 controllermanager.go:497] Started "statefulset"
I0330 19:57:08.769144 13668 controllermanager.go:497] Started "cronjob"
I0330 19:57:08.769572 13668 controllermanager.go:497] Started "csrapproving"
I0330 19:57:08.770288 13668 controllermanager.go:497] Started "persistentvolume-binder"
I0330 19:57:08.767921 13668 garbagecollector.go:130] Starting garbage collector controller
I0330 19:57:08.770473 13668 controller_utils.go:1027] Waiting for caches to sync for garbage collector
controller
I0330 19:57:08.770496 13668 ttl_controller.go:116] Starting TTL controller
I0330 19:57:08.770508 13668 controller_utils.go:1027] Waiting for caches to sync for TTL controller
I0330 19:57:08.770525 13668 stateful_set.go:151] Starting stateful set controller
I0330 19:57:08.770552 13668 controller_utils.go:1027] Waiting for caches to sync for stateful set controller
I0330 19:57:08.770569 13668 graph_builder.go:308] GraphBuilder running
I0330 19:57:08.770584 13668 cronjob_controller.go:94] Starting CronJob Manager
I0330 19:57:08.770747 13668 certificate_controller.go:113] Starting certificate controller
I0330 19:57:08.770766 13668 controller_utils.go:1027] Waiting for caches to sync for certificate controller
I0330 19:57:08.770787 13668 pv_controller_base.go:270] Starting persistent volume controller
I0330 19:57:08.770801 13668 controller_utils.go:1027] Waiting for caches to sync for persistent volume
controller
I0330 19:57:08.771101 13668 controllermanager.go:497] Started "pv-protection"
I0330 19:57:08.771228 13668 pv_protection_controller.go:81] Starting PV protection controller
I0330 19:57:08.771812 13668 controller_utils.go:1027] Waiting for caches to sync for PV protection controller

```



```

I0330 19:57:08.772111 13668 controllermanager.go:497] Started "podgc"
I0330 19:57:08.772254 13668 gc_controller.go:76] Starting GC controller
I0330 19:57:08.772272 13668 controller_utils.go:1027] Waiting for caches to sync for GC controller
I0330 19:57:08.772982 13668 controllermanager.go:497] Started "serviceaccount"
I0330 19:57:08.773006 13668 serviceaccounts_controller.go:115] Starting service account controller
I0330 19:57:08.773049 13668 controller_utils.go:1027] Waiting for caches to sync for service account controller
I0330 19:57:08.773249 13668 controllermanager.go:497] Started "csrcleaner"
W0330 19:57:08.773263 13668 controllermanager.go:489] Skipping "nodeipam"
I0330 19:57:08.773281 13668 cleaner.go:81] Starting CSR cleaner controller
I0330 19:57:08.775263 13668 controllermanager.go:497] Started "replicationcontroller"
I0330 19:57:08.776190 13668 controllermanager.go:497] Started "replicaset"
I0330 19:57:08.776217 13668 core.go:171] Will not configure cloud provider routes for allocate-node-cidrs:
false, configure-cloud-routes: true.
W0330 19:57:08.776229 13668 controllermanager.go:489] Skipping "route"
I0330 19:57:08.776848 13668 node_lifecycle_controller.go:77] Sending events to api server
E0330 19:57:08.776911 13668 core.go:161] failed to start cloud node lifecycle controller: no cloud provider
provided
W0330 19:57:08.776942 13668 controllermanager.go:489] Skipping "cloud-node-lifecycle"
I0330 19:57:08.778554 13668 controllermanager.go:497] Started "attachdetach"
I0330 19:57:08.780283 13668 replica_set.go:182] Starting replicaset controller
I0330 19:57:08.780311 13668 controller_utils.go:1027] Waiting for caches to sync for ReplicaSet controller
I0330 19:57:08.780429 13668 controllermanager.go:497] Started "endpoint"
I0330 19:57:08.780941 13668 controllermanager.go:497] Started "daemonset"
E0330 19:57:08.781348 13668 core.go:77] Failed to start service controller: WARNING: no cloud provider provided,
services of type LoadBalancer will fail
W0330 19:57:08.781367 13668 controllermanager.go:489] Skipping "service"
W0330 19:57:08.781409 13668 controllermanager.go:489] Skipping "root-ca-cert-publisher"
I0330 19:57:08.782137 13668 replica_set.go:182] Starting replicationcontroller controller
I0330 19:57:08.782189 13668 controller_utils.go:1027] Waiting for caches to sync for ReplicationController
controller
I0330 19:57:08.782240 13668 endpoints_controller.go:166] Starting endpoint controller
I0330 19:57:08.782278 13668 controller_utils.go:1027] Waiting for caches to sync for endpoint controller
I0330 19:57:08.782324 13668 daemon_controller.go:267] Starting daemon sets controller
I0330 19:57:08.782360 13668 controller_utils.go:1027] Waiting for caches to sync for daemon sets controller
I0330 19:57:08.786116 13668 attach_detach_controller.go:323] Starting attach detach controller
I0330 19:57:08.786131 13668 controller_utils.go:1027] Waiting for caches to sync for attach detach controller
I0330 19:57:08.983474 13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
endpoints
I0330 19:57:08.983556 13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
replicasets.apps
I0330 19:57:08.983598 13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
cronjobs.batch
I0330 19:57:08.983629 13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for

```

```

daemonsets.extensions
I0330 19:57:08.983667    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
replicasets.extensions
I0330 19:57:08.983708    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
horizontalpodautoscalers.autoscaling
I0330 19:57:08.983744    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
podd disruptionbudgets.policy
I0330 19:57:08.983777    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
leases.coordination.k8s.io
I0330 19:57:08.983896    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
ingresses.extensions
I0330 19:57:08.983949    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
controllerrevisions.apps
I0330 19:57:08.983987    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
events.events.k8s.io
I0330 19:57:08.984021    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
ingresses.networking.k8s.io
I0330 19:57:08.984080    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
limitranges
I0330 19:57:08.984145    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
daemonsets.apps
I0330 19:57:08.984234    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
statefulsets.apps
I0330 19:57:08.984281    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
networkpolicies.networking.k8s.io
I0330 19:57:08.984424    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
rolebindings.rbac.authorization.k8s.io
I0330 19:57:08.984475    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
serviceaccounts
I0330 19:57:08.984522    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
podtemplates
I0330 19:57:08.984600    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
deployments.extensions
I0330 19:57:08.984638    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
deployments.apps
I0330 19:57:08.984671    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
jobs.batch
I0330 19:57:08.984728    13668 resource_quota_monitor.go:228] QuotaMonitor created object count evaluator for
roles.rbac.authorization.k8s.io
E0330 19:57:08.984860    13668 resource_quota_controller.go:171] initial monitor sync has error: couldn't start
monitor for resource "extensions/v1beta1, Resource=networkpolicies": unable to monitor quota for resource
"extensions/v1beta1, Resource=networkpolicies"
I0330 19:57:08.984879    13668 controllermanager.go:497] Started "resourcequota"

```

```

I0330 19:57:08.985212 13668 resource_quota_controller.go:276] Starting resource quota controller
I0330 19:57:08.985234 13668 controller_utils.go:1027] Waiting for caches to sync for resource quota controller
I0330 19:57:09.000180 13668 resource_quota_monitor.go:301] QuotaMonitor running
W0330 19:57:09.008060 13668 actual_state_of_world.go:503] Failed to update statusUpdateNeeded field in actual
state of world: Failed to set statusUpdateNeeded to needed true, because nodeName="nodea" does not exist
I0330 19:57:09.037886 13668 controller_utils.go:1034] Caches are synced for HPA controller
I0330 19:57:09.038094 13668 controller_utils.go:1034] Caches are synced for job controller
I0330 19:57:09.061441 13668 controller_utils.go:1034] Caches are synced for namespace controller
I0330 19:57:09.070670 13668 controller_utils.go:1034] Caches are synced for TTL controller
I0330 19:57:09.070887 13668 controller_utils.go:1034] Caches are synced for certificate controller
I0330 19:57:09.072045 13668 controller_utils.go:1034] Caches are synced for PV protection controller
I0330 19:57:09.072381 13668 controller_utils.go:1034] Caches are synced for GC controller
I0330 19:57:09.073692 13668 controller_utils.go:1034] Caches are synced for service account controller
I0330 19:57:09.082410 13668 controller_utils.go:1034] Caches are synced for daemon sets controller
I0330 19:57:09.082503 13668 controller_utils.go:1034] Caches are synced for ReplicationController controller
I0330 19:57:09.086296 13668 controller_utils.go:1034] Caches are synced for attach detach controller
I0330 19:57:09.136943 13668 controller_utils.go:1034] Caches are synced for expand controller
I0330 19:57:09.137202 13668 controller_utils.go:1034] Caches are synced for PVC protection controller
I0330 19:57:09.162369 13668 controller_utils.go:1034] Caches are synced for ClusterRoleAggregator controller
I0330 19:57:09.437029 13668 controller_utils.go:1034] Caches are synced for taint controller
I0330 19:57:09.437092 13668 node_lifecycle_controller.go:1159] Initializing eviction metric for zone:
W0330 19:57:09.437143 13668 node_lifecycle_controller.go:833] Missing timestamp for Node nodea. Assuming now as
a timestamp.
I0330 19:57:09.437176 13668 node_lifecycle_controller.go:1059] Controller detected that zone is now in state
Normal.
I0330 19:57:09.437712 13668 taint_manager.go:198] Starting NoExecuteTaintManager
I0330 19:57:09.438138 13668 event.go:209] Event(v1.ObjectReference{Kind:"Node", Namespace:"", Name:"nodea",
UID:"f17942e8-5322-11e9-a6c5-02ef63d53bbe", APIVersion:"", ResourceVersion:"", FieldPath:"}): type: 'Normal'
reason: 'RegisteredNode' Node nodea event: Registered Node nodea in Controller
I0330 19:57:09.462156 13668 controller_utils.go:1034] Caches are synced for deployment controller
I0330 19:57:09.480489 13668 controller_utils.go:1034] Caches are synced for ReplicaSet controller
I0330 19:57:09.635860 13668 controller_utils.go:1034] Caches are synced for disruption controller
I0330 19:57:09.635884 13668 disruption.go:294] Sending events to api server.
I0330 19:57:09.670678 13668 controller_utils.go:1034] Caches are synced for stateful set controller
I0330 19:57:09.670964 13668 controller_utils.go:1034] Caches are synced for persistent volume controller
I0330 19:57:09.782420 13668 controller_utils.go:1034] Caches are synced for endpoint controller
I0330 19:57:09.785420 13668 controller_utils.go:1034] Caches are synced for resource quota controller
I0330 19:57:09.871514 13668 controller_utils.go:1034] Caches are synced for garbage collector controller
I0330 19:57:09.871537 13668 garbagecollector.go:139] Garbage collector: all resource monitors have synced.
Proceeding to collect garbage
I0330 19:57:10.263511 13668 controller_utils.go:1027] Waiting for caches to sync for garbage collector
controller
I0330 19:57:10.363817 13668 controller_utils.go:1034] Caches are synced for garbage collector controller

```

```
E0330 19:57:10.683383 13668 resource_quota_controller.go:437] failed to sync resource monitors: couldn't start monitor for resource "extensions/v1beta1, Resource=networkpolicies": unable to monitor quota for resource "extensions/v1beta1, Resource=networkpolicies"
```

This provides us with a lot of output.

A key takeaway:

```
I0330 19:57:08.122732 13668 leaderelection.go:227] successfully acquired lease kube-system/kube-controller-manager
```

Like the scheduler, there can only be one active controller manager in a cluster. Anytime a new controller manager starts it forces an election. Given that this is the first controller manager it becomes the leader and will actively begin managing Deployments, ReplicaSets and Replication Controllers.

The Controller Manager manages many other resources however. Look over the log output and identify the various resource types reported.

You should be able to find at least:

- ReplicationController
- DaemonSet
- Job
- Deployment
- ReplicaSet
- HorizontalPodAutoscaler
- StatefulSet

After running the controller manager, our pods should successfully terminate and give back control over our terminal:

```
pod "nginx" deleted
pod "nginx-23152-a" deleted
pod "nginx-25624-a" deleted
pod "nginx-25979-a" deleted
pod "nginx-29538-a" deleted
pod "nginx-9358-a" deleted
pod "nginx-a" deleted
pod "nginx-b" deleted
```

## 14. Running deployments without a controller manager

To get an even clearer picture of our cluster's function with and without a Controller Manager, let's build a test deployment and see what happens when we create it without a Controller Manager. Perform the following steps on nodea:

**Stop the controller manager** with Ctrl+C (^C):

```
...
I0330 20:00:20.223408    15191 controller_utils.go:1034] Caches are synced for attach detach controller
I0330 20:00:20.292498    15191 controller_utils.go:1034] Caches are synced for service account controller
I0330 20:00:20.320529    15191 controller_utils.go:1034] Caches are synced for namespace controller
I0330 20:00:20.406418    15191 controller_utils.go:1034] Caches are synced for ReplicationController controller
I0330 20:00:20.406613    15191 controller_utils.go:1034] Caches are synced for disruption controller
I0330 20:00:20.406693    15191 disruption.go:294] Sending events to api server.
I0330 20:00:20.697002    15191 controller_utils.go:1034] Caches are synced for certificate controller
I0330 20:00:20.796654    15191 controller_utils.go:1034] Caches are synced for garbage collector controller
I0330 20:00:20.796677    15191 garbagecollector.go:139] Garbage collector: all resource monitors have synced.
Proceeding to collect garbage
I0330 20:00:20.809309    15191 controller_utils.go:1034] Caches are synced for resource quota controller
I0330 20:00:21.192971    15191 controller_utils.go:1027] Waiting for caches to sync for garbage collector
controller
I0330 20:00:21.293201    15191 controller_utils.go:1034] Caches are synced for garbage collector controller
E0330 20:00:21.607165    15191 resource_quota_controller.go:437] failed to sync resource monitors: couldn't start
monitor for resource "extensions/v1beta1, Resource=networkpolicies": unable to monitor quota for resource
"extensions/v1beta1, Resource=networkpolicies"
```

```
^C
ubuntu@nodea:~$
```

Now build the test deployment:

```
ubuntu@nodea:~$ vim testdep.yaml
ubuntu@nodea:~$ cat testdep.yaml
```

```
apiVersion: apps/v1
kind: Deployment
```

```
metadata:
  name: nginx-deployment
  labels:
    name: nginx-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl create -f testdep.yaml

deployment.apps/nginx-deployment created
ubuntu@nodea:~$
```

Note that the message 'deployment "nginx-deployment" created' is from the API Server and indicates nothing other than that the API server added your desired state to the etcd store (this will fail in only the most dire circumstances).

Try listing the pods in your cluster:

```
ubuntu@nodea:~$ kubectl get pods

No resources found.
ubuntu@nodea:~$
```

This is a bad sign. Why didn't the cluster create the 2 pods requested?

List the other resources that should be associated with your deployment

```
ubuntu@nodea:~$ kubectl get rs
```

No resources found.

```
ubuntu@nodea:~$
```

```
ubuntu@nodea:~$ kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-deployment	0/2	0	0	18s

```
ubuntu@nodea:~$
```

So our deployment was added to the cluster target state but nothing else was. Let's look deeper:

```
ubuntu@nodea:~$ kubectl describe deploy nginx-deployment
```

```
Name: nginx-deployment
Namespace: default
CreationTimestamp: Sat, 30 Mar 2019 20:01:59 +0000
Labels: <none>
Annotations: <none>
Selector: app=nginx
Replicas: 2 desired | 0 updated | 0 total | 0 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels: app=nginx
  Containers:
    nginx:
      Image: nginx:1.7.9
      Port: 80/TCP
      Host Port: 0/TCP
      Environment: <none>
      Mounts: <none>
  Volumes: <none>
OldReplicaSets: <none>
```

```
NewReplicaSet:    <none>
Events:          <none>
ubuntu@nodea:~$
```

Our deployment has no ReplicaSets. As you have probably guessed this is because the Kubernetes component that acts on deployments and creates ReplicaSets is the Controller Manager and we stopped it.

Start Controller Manager once more:

```
ubuntu@nodea:~$ $HOME/k8s/_output/bin/kube-controller-manager --kubeconfig=nodea.conf
```

```
...
```

```
I0330 20:03:37.097757 15757 controller_utils.go:1034] Caches are synced for GC controller
I0330 20:03:37.099389 15757 controller_utils.go:1034] Caches are synced for PV protection controller
I0330 20:03:37.099570 15757 controller_utils.go:1034] Caches are synced for job controller
I0330 20:03:37.108021 15757 controller_utils.go:1034] Caches are synced for deployment controller
I0330 20:03:37.111620 15757 event.go:209] Event(v1.ObjectReference{Kind:"Deployment", Namespace:"default",
Name:"nginx-deployment", UID:"ad28fb15-5326-11e9-92a6-02ef63d53bbe", APIVersion:"apps/v1", ResourceVersion:"1225",
FieldPath:""}): type: 'Normal' reason: 'ScalingReplicaSet' Scaled up replica set nginx-deployment-6dd86d77d to 2
I0330 20:03:37.115271 15757 controller_utils.go:1034] Caches are synced for ClusterRoleAggregator controller
I0330 20:03:37.130847 15757 controller_utils.go:1034] Caches are synced for namespace controller
I0330 20:03:37.131296 15757 controller_utils.go:1034] Caches are synced for service account controller
I0330 20:03:37.131476 15757 controller_utils.go:1034] Caches are synced for ReplicaSet controller
I0330 20:03:37.134585 15757 controller_utils.go:1034] Caches are synced for certificate controller
I0330 20:03:37.134616 15757 controller_utils.go:1034] Caches are synced for disruption controller
I0330 20:03:37.134651 15757 disruption.go:294] Sending events to api server.
I0330 20:03:37.134938 15757 event.go:209] Event(v1.ObjectReference{Kind:"ReplicaSet", Namespace:"default",
Name:"nginx-deployment-6dd86d77d", UID:"e79b48d1-5326-11e9-92a6-02ef63d53bbe", APIVersion:"apps/v1",
ResourceVersion:"1299", FieldPath:""}): type: 'Normal' reason: 'SuccessfulCreate' Created pod: nginx-deployment-
6dd86d77d-kcsdc
I0330 20:03:37.137283 15757 event.go:209] Event(v1.ObjectReference{Kind:"ReplicaSet", Namespace:"default",
Name:"nginx-deployment-6dd86d77d", UID:"e79b48d1-5326-11e9-92a6-02ef63d53bbe", APIVersion:"apps/v1",
ResourceVersion:"1299", FieldPath:""}): type: 'Normal' reason: 'SuccessfulCreate' Created pod: nginx-deployment-
6dd86d77d-z6tdf
I0330 20:03:37.147939 15757 controller_utils.go:1034] Caches are synced for expand controller
I0330 20:03:37.148035 15757 controller_utils.go:1034] Caches are synced for endpoint controller
I0330 20:03:37.148098 15757 controller_utils.go:1034] Caches are synced for PVC protection controller
I0330 20:03:37.148746 15757 controller_utils.go:1034] Caches are synced for stateful set controller
I0330 20:03:37.150774 15757 controller_utils.go:1034] Caches are synced for ReplicationController controller
I0330 20:03:37.510259 15757 controller_utils.go:1034] Caches are synced for HPA controller
```



```

W0330 20:03:37.803298 15757 actual_state_of_world.go:503] Failed to update statusUpdateNeeded field in actual
state of world: Failed to set statusUpdateNeeded to needed true, because nodeName="nodea" does not exist
I0330 20:03:37.848329 15757 controller_utils.go:1034] Caches are synced for TTL controller
I0330 20:03:37.849448 15757 controller_utils.go:1034] Caches are synced for taint controller
I0330 20:03:37.849502 15757 node_lifecycle_controller.go:1159] Initializing eviction metric for zone:
W0330 20:03:37.849556 15757 node_lifecycle_controller.go:833] Missing timestamp for Node nodea. Assuming now as
a timestamp.
I0330 20:03:37.849598 15757 node_lifecycle_controller.go:1059] Controller detected that zone is now in state
Normal.
I0330 20:03:37.849631 15757 event.go:209] Event(v1.ObjectReference{Kind:"Node", Namespace:"", Name:"nodea",
UID:"f17942e8-5322-11e9-a6c5-02ef63d53bbe", APIVersion:"", ResourceVersion:"", FieldPath:""}): type: 'Normal'
reason: 'RegisteredNode' Node nodea event: Registered Node nodea in Controller
I0330 20:03:37.849655 15757 taint_manager.go:198] Starting NoExecuteTaintManager
I0330 20:03:37.896796 15757 controller_utils.go:1034] Caches are synced for daemon sets controller
I0330 20:03:37.898246 15757 controller_utils.go:1034] Caches are synced for persistent volume controller
I0330 20:03:37.899307 15757 controller_utils.go:1034] Caches are synced for attach detach controller
I0330 20:03:37.919088 15757 controller_utils.go:1034] Caches are synced for garbage collector controller
I0330 20:03:37.919108 15757 garbagecollector.go:139] Garbage collector: all resource monitors have synced.
Proceeding to collect garbage
I0330 20:03:37.947852 15757 controller_utils.go:1034] Caches are synced for resource quota controller
I0330 20:03:38.316138 15757 controller_utils.go:1027] Waiting for caches to sync for garbage collector
controller
I0330 20:03:38.416371 15757 controller_utils.go:1034] Caches are synced for garbage collector controller
E0330 20:03:38.746731 15757 resource_quota_controller.go:437] failed to sync resource monitors: couldn't start
monitor for resource "extensions/v1beta1, Resource=networkpolicies": unable to monitor quota for resource
"extensions/v1beta1, Resource=networkpolicies"

```

Once the Controller Manager is up and running, toward the end of the log output you will see it discover your deployment.

```

I0330 20:03:37.111620 15757 event.go:209] Event(v1.ObjectReference{Kind:"Deployment", Namespace:"default", Name:"nginx-deployment",
UID:"ad28fb15-5326-11e9-92a6-02ef63d53bbe", APIVersion:"apps/v1", ResourceVersion:"1225", FieldPath:""}): type: 'Normal' reason:
'ScalingReplicaSet' Scaled up replica set nginx-deployment-6dd86d77d to 2

```

This is immediately followed by events reporting the actions take by the Controller Manager to bring the cluster in line with your wishes:

```

I0330 20:03:37.134938 15757 event.go:209] Event(v1.ObjectReference{Kind:"ReplicaSet", Namespace:"default", Name:"nginx-deployment-
6dd86d77d", UID:"e79b48d1-5326-11e9-92a6-02ef63d53bbe", APIVersion:"apps/v1", ResourceVersion:"1299", FieldPath:""}): type: 'Normal' reason:
'SuccessfulCreate' Created pod: nginx-deployment-6dd86d77d-kcsdc

```

```
I0330 20:03:37.137283 15757 event.go:209] Event(v1.ObjectReference{Kind:"ReplicaSet", Namespace:"default", Name:"nginx-deployment-6dd86d77d", UID:"e79b48d1-5326-11e9-92a6-02ef63d53bbe", APIVersion:"apps/v1", ResourceVersion:"1299", FieldPath:""}): type: 'Normal' reason: 'SuccessfulCreate' Created pod: nginx-deployment-6dd86d77d-z6tdf
```

Now try displaying the active pods:

```
ubuntu@nodea:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-6dd86d77d-jxxx7	1/1	Running	0	101s
nginx-deployment-6dd86d77d-qz8nl	1/1	Running	0	101s

```
ubuntu@nodea:~$
```

As advertised, the Controller manager has created the two pods required. You can examine the system events to see the progression of work involved in launching your two pods:

```
ubuntu@nodea:~$ kubectl get events | grep $(kubectl get pod -o name |tail -1)
```

3m49s	Normal	Scheduled	pod/nginx-deployment-6dd86d77d-qz8nl	Successfully assigned
default/nginx-deployment-6dd86d77d-qz8nl to nodea				
6s	Warning	MissingClusterDNS	pod/nginx-deployment-6dd86d77d-qz8nl	pod: "nginx-deployment-6dd86d77d-qz8nl_default(231337e4-5328-11e9-92a6-02ef63d53bbe)". kubelet does not have ClusterDNS IP configured and cannot create Pod using "ClusterFirst" policy. Falling back to "Default" policy.
3m48s	Normal	Pulled	pod/nginx-deployment-6dd86d77d-qz8nl	Container image
"nginx:1.7.9" already present on machine				
3m48s	Normal	Created	pod/nginx-deployment-6dd86d77d-qz8nl	Created container nginx
3m48s	Normal	Started	pod/nginx-deployment-6dd86d77d-qz8nl	Started container nginx

```
ubuntu@nodea:~$
```

After creating the Deployment's ReplicaSet the Controller Manager creates two Pods from the template and submits them to the cluster. The Scheduler is exclusively responsible for scheduling Pods to Kubelets when the Pods are not pinned to a particular host.

As soon as the Pods are scheduled the Kubelets begin pulling images and launching Pods. When the Pods are up the ReplicaSet reports successful Pod creation and when all Pods are up the Deployment reports successful scaling (from 0 to 2 as the Controller Manager logs stated).

Note that all Pods are automatically registered in the Kubernetes Cluster DNS if one is configured. We see Warnings because we have not yet setup Cluster DNS. We will take care of that in a future lab.

Relist your running Pods, ReplicaSets, and Deployments:

```
ubuntu@nodea:~$ kubectl get deploy,rs,po
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.extensions/nginx-deployment	2/2	2	2	6m29s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.extensions/nginx-deployment-6dd86d77d	2	2	2	6m29s

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-deployment-6dd86d77d-jxxx7	1/1	Running	0	6m29s
pod/nginx-deployment-6dd86d77d-qz8nl	1/1	Running	0	6m29s

```
ubuntu@nodea:~$
```

Everything looks healthy! Delete your deployment and all pods.

## [OPTIONAL] Customizing the scheduler

The Kubernetes scheduler makes use of predicates to identify nodes to which a pod may be scheduled. The scheduler then uses policies to rank the nodes that pass the predicate tests. By default, Kubernetes provides built-in predicates and priority policies documented in [scheduler\\_algorithm.md](#). The predicates and priorities code are defined in [plugin/pkg/scheduler/algorithm/predicates/predicates.go](#) and [plugin/pkg/scheduler/algorithm/priorities](#), respectively.

The policies that are applied when scheduling can be selected in one of two ways. The default policies used are selected by the functions [defaultPredicates\(\)](#) and [defaultPriorities\(\)](#) in [plugin/pkg/scheduler/algorithmprovider/defaults/defaults.go](#). However, the choice of policies can be overridden by passing the command-line flag [--policy-config-file](#) to the scheduler, pointing to a JSON file specifying which scheduling policies to use. See [examples/scheduler-policy-config.json](#) for an example config file. Note that the config file format is versioned; the API is defined in [plugin/pkg/scheduler/api](#). To add a new scheduling policy, you should modify [plugin/pkg/scheduler/algorithm/predicates/predicates.go](#) or add to the directory [plugin/pkg/scheduler/algorithm/priorities](#), and either register the policy in [defaultPredicates\(\)](#) or [defaultPriorities\(\)](#), or use a policy config file.

To experiment with scheduling configuration, create a custom policy file and restart the scheduler with it.

```
ubuntu@nodea:~$ vim custom.json
ubuntu@nodea:~$ cat custom.json
```

```
{
"kind" : "Policy",
"apiVersion" : "v1",
"predicates" : [
  {"name" : "PodFitsHostPorts", "order": 1},
  {"name" : "PodFitsResources", "order": 3},
  {"name" : "NoDiskConflict", "order": 4},
  {"name" : "NoVolumeZoneConflict", "order": 6},
  {"name" : "MatchNodeSelector", "order": 5},
  {"name" : "HostName", "order": 2}
],
"priorities" : [
  {"name" : "LeastRequestedPriority", "weight" : 1},
  {"name" : "BalancedResourceAllocation", "weight" : 1},
  {"name" : "ServiceSpreadingPriority", "weight" : 1},
  {"name" : "EqualPriority", "weight" : 1}
]
}
```

```
ubuntu@nodea:~$
```

Now restart the scheduler with the new policy:

```
ubuntu@nodea:~$ $HOME/k8s/_output/bin/kube-scheduler --kubeconfig=nodea.conf \
--policy-config-file=custom.json --v=2

I0330 20:23:05.568546 17416 flags.go:33] FLAG: --address="0.0.0.0"
I0330 20:23:05.568605 17416 flags.go:33] FLAG: --algorithm-provider=""
I0330 20:23:05.568611 17416 flags.go:33] FLAG: --alsologtostderr="false"
I0330 20:23:05.568616 17416 flags.go:33] FLAG: --authentication-kubeconfig=""
I0330 20:23:05.568640 17416 flags.go:33] FLAG: --authentication-skip-lookup="false"
I0330 20:23:05.568647 17416 flags.go:33] FLAG: --authentication-token-webhook-cache-ttl="10s"
I0330 20:23:05.568653 17416 flags.go:33] FLAG: --authentication-tolerate-lookup-failure="true"
I0330 20:23:05.568658 17416 flags.go:33] FLAG: --authorization-always-allow-paths="/healthz"
```

```

I0330 20:23:05.568672 17416 flags.go:33] FLAG: --authorization-kubeconfig=""
I0330 20:23:05.568680 17416 flags.go:33] FLAG: --authorization-webhook-cache-authorized-ttl="10s"
I0330 20:23:05.568688 17416 flags.go:33] FLAG: --authorization-webhook-cache-unauthorized-ttl="10s"
I0330 20:23:05.568697 17416 flags.go:33] FLAG: --bind-address="0.0.0.0"
I0330 20:23:05.568706 17416 flags.go:33] FLAG: --cert-dir=""
I0330 20:23:05.568714 17416 flags.go:33] FLAG: --client-ca-file=""
I0330 20:23:05.568722 17416 flags.go:33] FLAG: --config=""
I0330 20:23:05.568730 17416 flags.go:33] FLAG: --contention-profiling="false"
I0330 20:23:05.568738 17416 flags.go:33] FLAG: --failure-domains="kubernetes.io/hostname,failure-
domain.beta.kubernetes.io/zone,failure-domain.beta.kubernetes.io/region"
I0330 20:23:05.568750 17416 flags.go:33] FLAG: --feature-gates=""
I0330 20:23:05.568760 17416 flags.go:33] FLAG: --hard-pod-affinity-symmetric-weight="1"
I0330 20:23:05.568765 17416 flags.go:33] FLAG: --help="false"
I0330 20:23:05.568769 17416 flags.go:33] FLAG: --http2-max-streams-per-connection="0"
I0330 20:23:05.568774 17416 flags.go:33] FLAG: --kube-api-burst="100"
I0330 20:23:05.568779 17416 flags.go:33] FLAG: --kube-api-content-type="application/vnd.kubernetes.protobuf"
I0330 20:23:05.568783 17416 flags.go:33] FLAG: --kube-api-qps="50"
I0330 20:23:05.568789 17416 flags.go:33] FLAG: --kubeconfig="nodea.conf"
I0330 20:23:05.568797 17416 flags.go:33] FLAG: --leader-elect="true"
I0330 20:23:05.568806 17416 flags.go:33] FLAG: --leader-elect-lease-duration="15s"
I0330 20:23:05.568814 17416 flags.go:33] FLAG: --leader-elect-renew-deadline="10s"
I0330 20:23:05.568822 17416 flags.go:33] FLAG: --leader-elect-resource-lock="endpoints"
I0330 20:23:05.568830 17416 flags.go:33] FLAG: --leader-elect-retry-period="2s"
I0330 20:23:05.568834 17416 flags.go:33] FLAG: --lock-object-name="kube-scheduler"
I0330 20:23:05.568842 17416 flags.go:33] FLAG: --lock-object-namespace="kube-system"
I0330 20:23:05.568878 17416 flags.go:33] FLAG: --log-backtrace-at=":0"
I0330 20:23:05.568890 17416 flags.go:33] FLAG: --log-dir=""
I0330 20:23:05.568898 17416 flags.go:33] FLAG: --log-file=""
I0330 20:23:05.568906 17416 flags.go:33] FLAG: --log-flush-frequency="5s"
I0330 20:23:05.568914 17416 flags.go:33] FLAG: --logtostderr="true"
I0330 20:23:05.568923 17416 flags.go:33] FLAG: --master=""
I0330 20:23:05.568927 17416 flags.go:33] FLAG: --policy-config-file="custom.json"
I0330 20:23:05.568935 17416 flags.go:33] FLAG: --policy-configmap=""
I0330 20:23:05.568943 17416 flags.go:33] FLAG: --policy-configmap-namespace="kube-system"
I0330 20:23:05.568950 17416 flags.go:33] FLAG: --port="10251"
I0330 20:23:05.568955 17416 flags.go:33] FLAG: --profiling="false"
I0330 20:23:05.568963 17416 flags.go:33] FLAG: --requestheader-allowed-names="[]"
I0330 20:23:05.568988 17416 flags.go:33] FLAG: --requestheader-client-ca-file=""
I0330 20:23:05.568996 17416 flags.go:33] FLAG: --requestheader-extra-headers-prefix="[x-remote-extra-]"
I0330 20:23:05.569008 17416 flags.go:33] FLAG: --requestheader-group-headers="[x-remote-group]"
I0330 20:23:05.569017 17416 flags.go:33] FLAG: --requestheader-username-headers="[x-remote-user]"
I0330 20:23:05.569027 17416 flags.go:33] FLAG: --scheduler-name="default-scheduler"
I0330 20:23:05.569039 17416 flags.go:33] FLAG: --secure-port="10259"

```

```

I0330 20:23:05.569045 17416 flags.go:33] FLAG: --skip-headers="false"
I0330 20:23:05.569053 17416 flags.go:33] FLAG: --stderrthreshold="2"
I0330 20:23:05.569075 17416 flags.go:33] FLAG: --tls-cert-file=""
I0330 20:23:05.569083 17416 flags.go:33] FLAG: --tls-cipher-suites="[]"
I0330 20:23:05.569088 17416 flags.go:33] FLAG: --tls-min-version=""
I0330 20:23:05.569096 17416 flags.go:33] FLAG: --tls-private-key-file=""
I0330 20:23:05.569103 17416 flags.go:33] FLAG: --tls-sni-cert-key="[]"
I0330 20:23:05.569113 17416 flags.go:33] FLAG: --use-legacy-policy-config="false"
I0330 20:23:05.569121 17416 flags.go:33] FLAG: --v="2"
I0330 20:23:05.569141 17416 flags.go:33] FLAG: --version="false"
I0330 20:23:05.569152 17416 flags.go:33] FLAG: --vmodule=""
I0330 20:23:05.569160 17416 flags.go:33] FLAG: --write-config-to=""
I0330 20:23:06.083603 17416 serving.go:319] Generated self-signed cert in-memory
W0330 20:23:06.726996 17416 authentication.go:249] No authentication-kubeconfig provided in order to lookup
client-ca-file in configmap/extension-apiserver-authentication in kube-system, so client certificate
authentication won't work.
W0330 20:23:06.727022 17416 authentication.go:252] No authentication-kubeconfig provided in order to lookup
requestheader-client-ca-file in configmap/extension-apiserver-authentication in kube-system, so request-header
client certificate authentication won't work.
W0330 20:23:06.727034 17416 authorization.go:146] No authorization-kubeconfig provided, so SubjectAccessReview
of authorization tokens won't work.
I0330 20:23:06.728618 17416 server.go:142] Version: v1.14.0
I0330 20:23:06.728663 17416 defaults.go:87] TaintNodesByCondition is enabled, PodToleratesNodeTaints predicate
is mandatory
I0330 20:23:06.728679 17416 server.go:161] Starting Kubernetes Scheduler version v1.14.0
I0330 20:23:06.729579 17416 factory.go:341] Creating scheduler from configuration: {{ } [{PodFitsHostPorts
<nil>} {PodFitsResources <nil>} {NoDiskConflict <nil>} {NoVolumeZoneConflict <nil>} {MatchNodeSelector <nil>}
{HostName <nil>}] [{LeastRequestedPriority 1 <nil>} {BalancedResourceAllocation 1 <nil>} {ServiceSpreadingPriority
1 <nil>} {EqualPriority 1 <nil>}]} [] 0 false}
I0330 20:23:06.729631 17416 factory.go:358] Registering predicate: PodFitsHostPorts
I0330 20:23:06.729648 17416 plugins.go:236] Predicate type PodFitsHostPorts already registered, reusing.
I0330 20:23:06.729659 17416 factory.go:358] Registering predicate: PodFitsResources
I0330 20:23:06.729667 17416 plugins.go:236] Predicate type PodFitsResources already registered, reusing.
I0330 20:23:06.729676 17416 factory.go:358] Registering predicate: NoDiskConflict
I0330 20:23:06.729684 17416 plugins.go:236] Predicate type NoDiskConflict already registered, reusing.
I0330 20:23:06.729692 17416 factory.go:358] Registering predicate: NoVolumeZoneConflict
I0330 20:23:06.729701 17416 plugins.go:236] Predicate type NoVolumeZoneConflict already registered, reusing.
I0330 20:23:06.729709 17416 factory.go:358] Registering predicate: MatchNodeSelector
I0330 20:23:06.729726 17416 plugins.go:236] Predicate type MatchNodeSelector already registered, reusing.
I0330 20:23:06.729734 17416 factory.go:358] Registering predicate: HostName
I0330 20:23:06.729743 17416 plugins.go:236] Predicate type HostName already registered, reusing.
I0330 20:23:06.729757 17416 factory.go:373] Registering priority: LeastRequestedPriority
I0330 20:23:06.729767 17416 plugins.go:348] Priority type LeastRequestedPriority already registered, reusing.

```

```

I0330 20:23:06.729779    17416 factory.go:373] Registering priority: BalancedResourceAllocation
I0330 20:23:06.729788    17416 plugins.go:348] Priority type BalancedResourceAllocation already registered,
reusing.
I0330 20:23:06.729797    17416 factory.go:373] Registering priority: ServiceSpreadingPriority
I0330 20:23:06.729806    17416 plugins.go:348] Priority type ServiceSpreadingPriority already registered, reusing.
I0330 20:23:06.729816    17416 factory.go:373] Registering priority: EqualPriority
I0330 20:23:06.729824    17416 plugins.go:348] Priority type EqualPriority already registered, reusing.
I0330 20:23:06.729835    17416 factory.go:412] Creating scheduler with fit predicates 'map[HostName:{}
MatchNodeSelector:{} NoDiskConflict:{} NoVolumeZoneConflict:{} PodFitsHostPorts:{} PodFitsResources:{}]' and
priority functions 'map[BalancedResourceAllocation:{} EqualPriority:{} LeastRequestedPriority:{}
ServiceSpreadingPriority:{}]'
W0330 20:23:06.730501    17416 authorization.go:47] Authorization is disabled
W0330 20:23:06.730517    17416 authentication.go:55] Authentication is disabled
I0330 20:23:06.730530    17416 deprecated_insecure_serving.go:49] Serving healthz insecurely on [::]:10251
I0330 20:23:06.731220    17416 secure_serving.go:116] Serving securely on [::]:10259
I0330 20:23:07.633103    17416 controller_utils.go:1027] Waiting for caches to sync for scheduler controller
I0330 20:23:07.733332    17416 controller_utils.go:1034] Caches are synced for scheduler controller
I0330 20:23:07.733406    17416 leaderelection.go:217] attempting to acquire leader lease  kube-system/kube-
scheduler...
I0330 20:23:25.141662    17416 leaderelection.go:227] successfully acquired lease kube-system/kube-scheduler

...

```

Launch several pods and adjust the values above to see how your changes effect pod placement.

```

ubuntu@nodea:~$ sed -e '/nodeName/d' testpod-a.yaml \
-e "s/name: nginx/name: nginx-$RANDOM/g" | kubectl create -f -

pod/nginx-708-a created
ubuntu@nodea:~$

```

```

ubuntu@nodea:~$ kubectl get events -w --sort-by=LASTSEEN

25m          Normal    RegisteredNode          node/nodea          Node nodea event:
Registered Node nodea in Controller
22m          Normal    RegisteredNode          node/nodea          Node nodea event:
Registered Node nodea in Controller
0s           Warning   MissingClusterDNS       pod/nginx-18049-a   pod: "nginx-18049-
a_default(fbdd1e4c-5329-11e9-92a6-02ef63d53bbe)". kubelet does not have ClusterDNS IP configured and cannot create

```

```

Pod using "ClusterFirst" policy. Falling back to "Default" policy.
0s      Warning   MissingClusterDNS      pod/nginx-2616-a      pod: "nginx-2616-
a_default(f9a6deb0-5329-11e9-92a6-02ef63d53bbe)". kubelet does not have ClusterDNS IP configured and cannot create
Pod using "ClusterFirst" policy. Falling back to "Default" policy.
0s      Warning   MissingClusterDNS      node/nodea            kubelet does not have
ClusterDNS IP configured and cannot create Pod using "ClusterFirst" policy. Falling back to "Default" policy.
0s      Warning   MissingClusterDNS      pod/nginx-708-a      pod: "nginx-708-
a_default(f1d83762-5329-11e9-92a6-02ef63d53bbe)". kubelet does not have ClusterDNS IP configured and cannot create
Pod using "ClusterFirst" policy. Falling back to "Default" policy.
0s      Warning   MissingClusterDNS      pod/nginx-18049-a     pod: "nginx-18049-
a_default(fbdd1e4c-5329-11e9-92a6-02ef63d53bbe)". kubelet does not have ClusterDNS IP configured and cannot create
Pod using "ClusterFirst" policy. Falling back to "Default" policy.
^[0s      Warning   MissingClusterDNS      pod/nginx-26007-a     pod: "nginx-26007-
a_default(fa3ba3a2-5329-11e9-92a6-02ef63d53bbe)". kubelet does not have ClusterDNS IP configured and cannot create
Pod using "ClusterFirst" policy. Falling back to "Default" policy.

...

```

You can review the functions here:

- <https://github.com/kubernetes/kubernetes/blob/master/pkg/scheduler/algorithm/predicates/predicates.go>

and here:

- <https://github.com/kubernetes/kubernetes/tree/master/pkg/scheduler/algorithm/priorities>

Delete all your existing pods and related resources but leave your kubelets, Scheduler, API Server and etcd running.

```

ubuntu@nodea:~$ kubectl get pod -o go-template \
--template '{{range .items}}{{.metadata.name}}{{"\n"}}{{end}}' | xargs kubectl delete pod

pod "nginx-13119" deleted
pod "nginx-16863" deleted
pod "nginx-17870" deleted
pod "nginx-32247" deleted
pod "nginx-5157" deleted

ubuntu@nodea:~$

```



Congratulations you have successfully completed the lab!

*Copyright (c) 2013-2019 RX-M LLC, Cloud Native Consulting, all rights reserved*