

# Advanced Kubernetes

## Lab 5 – Software Defined Networking (SDN)

Flannel is a virtual network that gives a subnet to each host for use with container runtimes.

Platforms like Google's Kubernetes assume that each container (pod) has a unique, routable IP inside the cluster. The advantage of this model is that it reduces the complexity of doing port mapping.

Before proceeding we will need to **stop all the cluster services**, then add the external interface to etcd and restart it. First retrieve your node's external IP and then configure etcd.

```
ubuntu@nodea:~$ ip a show eth0

2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc mq state UP group default qlen 1000
    link/ether 02:ef:63:d5:3b:be brd ff:ff:ff:ff:ff:ff
    inet 172.31.28.198/20 brd 172.31.31.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::ef:63ff:fed5:3bbe/64 scope link
        valid_lft forever preferred_lft forever

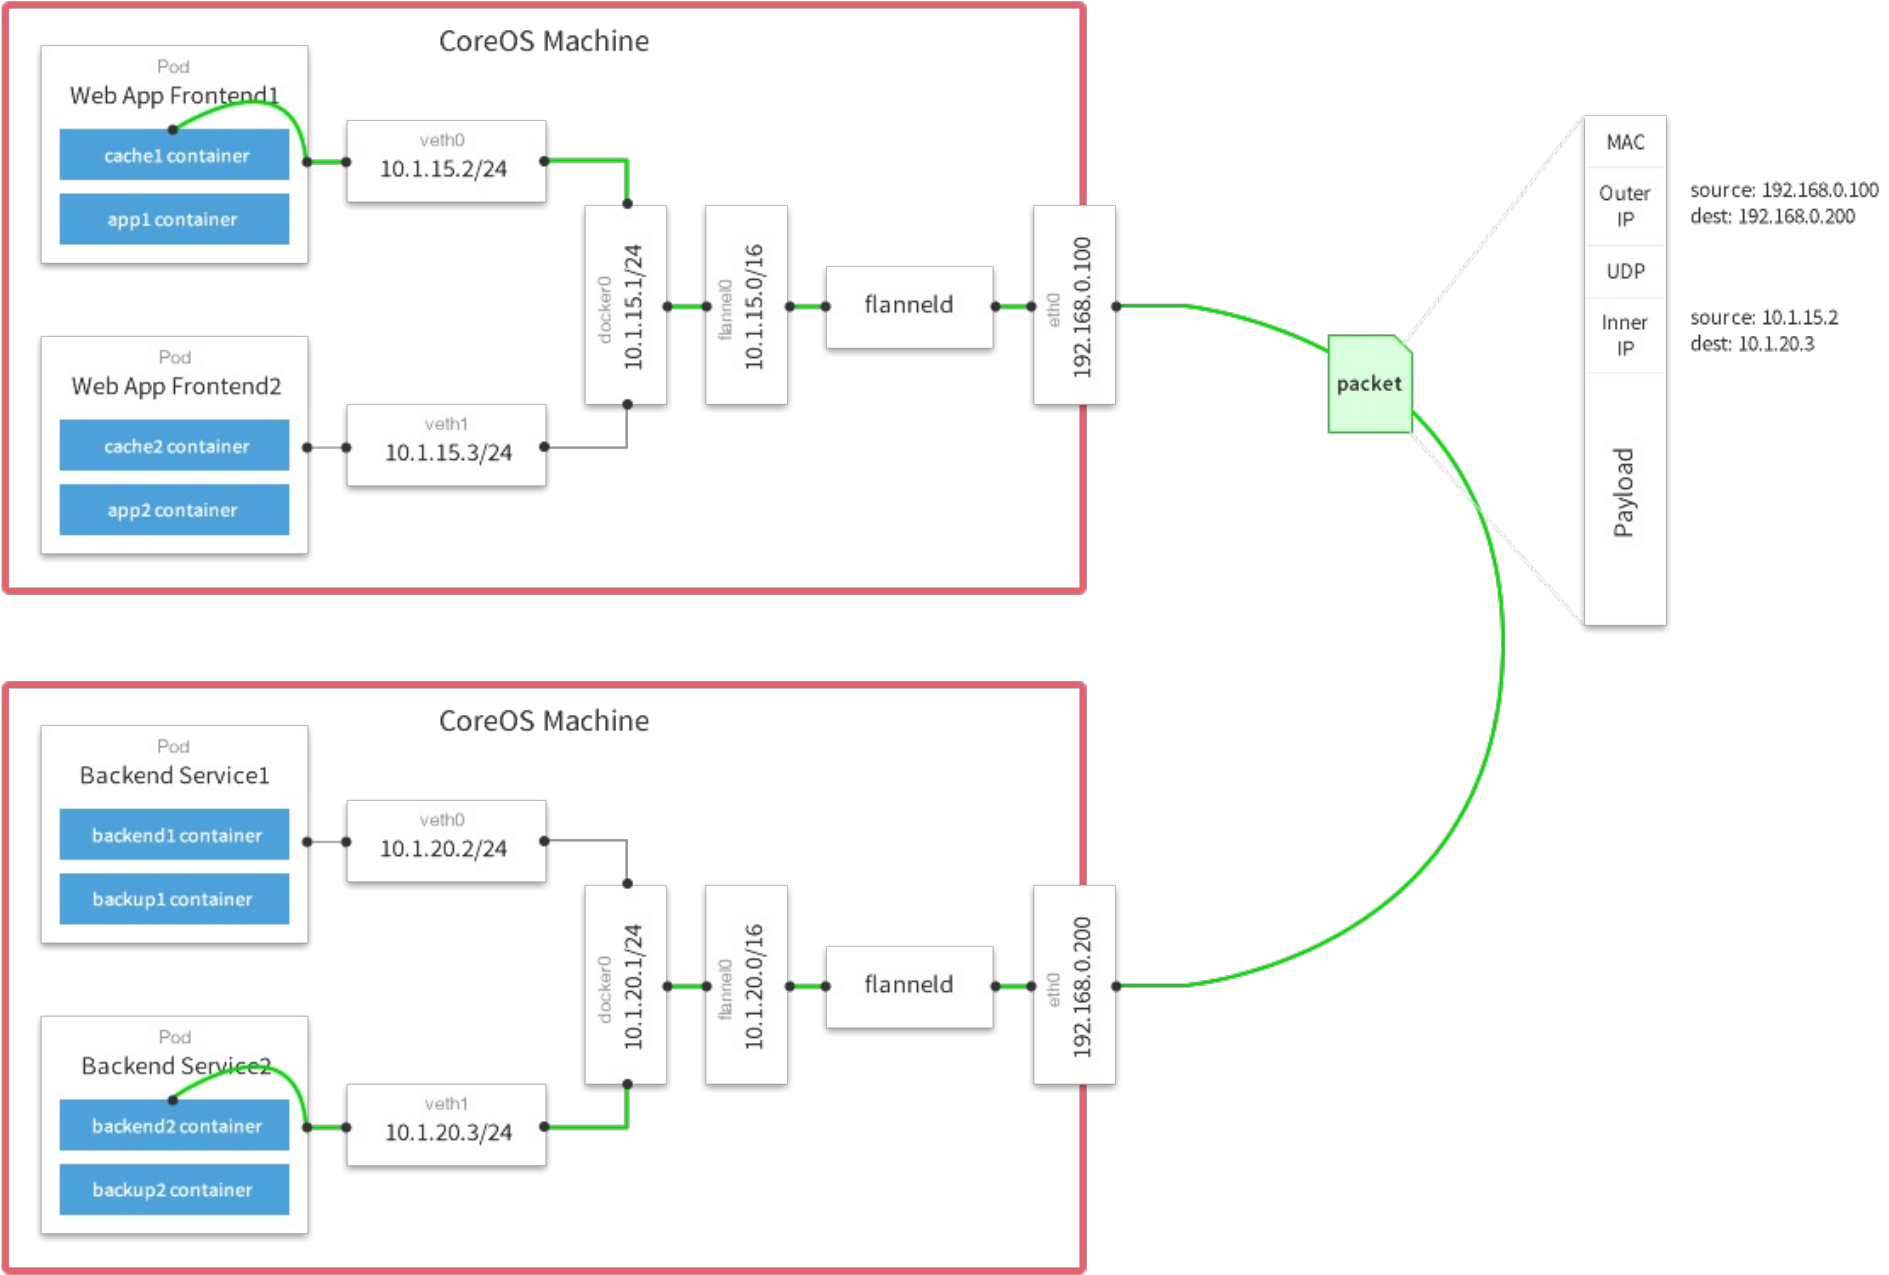
ubuntu@nodea:~$ rm -Rf ~/default.etcd/

ubuntu@nodea:~$ etcd --listen-client-urls 'http://0.0.0.0:2379' --advertise-client-urls
'http://172.31.28.198:2379'
```

### 1. flannel

flannel runs an agent, flanneld, on each host and is responsible for allocating a subnet lease out of a preconfigured address space. Flannel uses etcd to store the network configuration, allocated subnets, and auxiliary data (such as host's IP.) The forwarding of packets is achieved using one of several strategies that are known as backends. The simplest backend is UDP and uses a TUN device to encapsulate every IP fragment in a UDP packet, forming an overlay network. The

following diagram demonstrates the path a packet takes as it traverses the overlay network:



Flannel can also plug into Amazon VPC and use AWS routing tables to avoid the encapsulation associated with tunneling, which is how we will use it in this lab. You can find the latest details about flannel here <https://github.com/coreos/flannel>.

## Building flannel

On *nodea*

```
ubuntu@nodea:~$ cd $HOME

ubuntu@nodea:~$ export GOPATH=$HOME/goProjects/ && echo $GOPATH

/home/ubuntu/goProjects/

ubuntu@nodea:~$ go get -v github.com/coreos/flannel

github.com/coreos/flannel (download)
created GOPATH=/home/ubuntu/goProjects/; see 'go help gopath'
github.com/coreos/flannel/vendor/github.com/vishvananda/netns
github.com/coreos/flannel/vendor/github.com/golang/glog
github.com/coreos/flannel/vendor/github.com/vishvananda/netlink/nl
github.com/coreos/flannel/vendor/golang.org/x/net/context
github.com/coreos/flannel/vendor/github.com/denverdino/aliyungo/util
github.com/coreos/flannel/vendor/github.com/vishvananda/netlink
github.com/coreos/flannel/vendor/github.com/denverdino/aliyungo/common
github.com/coreos/flannel/vendor/github.com/denverdino/aliyungo/ecs

...

ubuntu@nodea:~$ cd $GOPATH/src/github.com/coreos/flannel

ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ go build -v github.com/coreos/flannel

ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ make dist/flanneld

go build -o dist/flanneld \
  -ldflags '-s -w -X github.com/coreos/flannel/version.Version=v0.11.0-11-ge4deb05 -extldflags "-static"'
# github.com/coreos/flannel
/tmp/go-link-934176633/000022.o: In function `mygetgrouplist':
/workdir/go/src/os/user/getgrouplist_unix.go:16: warning: Using 'getgrouplist' in statically linked applications
requires at runtime the shared libraries from the glibc version used for linking
/tmp/go-link-934176633/000021.o: In function `mygetgrgid_r':
```

```
/workdir/go/src/os/user/cgo_lookup_unix.go:38: warning: Using 'getgrgid_r' in statically linked applications
requires at runtime the shared libraries from the glibc version used for linking
/tmp/go-link-934176633/000021.o: In function `mygetgrnam_r':
/workdir/go/src/os/user/cgo_lookup_unix.go:43: warning: Using 'getgrnam_r' in statically linked applications
requires at runtime the shared libraries from the glibc version used for linking
/tmp/go-link-934176633/000021.o: In function `mygetpwnam_r':
/workdir/go/src/os/user/cgo_lookup_unix.go:33: warning: Using 'getpwnam_r' in statically linked applications
requires at runtime the shared libraries from the glibc version used for linking
/tmp/go-link-934176633/000021.o: In function `mygetpwuid_r':
/workdir/go/src/os/user/cgo_lookup_unix.go:28: warning: Using 'getpwuid_r' in statically linked applications
requires at runtime the shared libraries from the glibc version used for linking
/tmp/go-link-934176633/000004.o: In function `_cgo_7e1b3c2abc8d_C2func_getaddrinfo':
/tmp/go-build/cgo-gcc-prolog:57: warning: Using 'getaddrinfo' in statically linked applications requires at
runtime the shared libraries from the glibc version used for linking
```

```
ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ ./dist/flannel -version
```

```
v0.11.0-11-ge4deb05
```

```
ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$
```

Configure the subnet for our overlay; first, check what configuration, if any, is stored in etcd:

```
ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ etcdctl get /coreos.com/network/config
```

```
Error: 100: Key not found (/coreos.com) [3]
```

Now, set a network config with the type "aws-vpc". NOTE: the 10.1.0.0/16 network in the example below is for **example purposes ONLY**; your instructor will assign you a unique subnet to avoid overlapping with other students.

```
ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ etcdctl set /coreos.com/network/config \
'{"Network": "10.1.0.0/16", "Backend": {"Type": "aws-vpc"}}'
```

```
{"Network": "10.1.0.0/16", "Backend": {"Type": "aws-vpc"}}
```

Make sure the configuration took:

```
ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ etcdctl get /coreos.com/network/config

{"Network":"10.1.0.0/16", "Backend": {"Type": "aws-vpc"}}
ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$
```

## Starting flanneld

**Open a new terminal** and start the flanneld daemon, on *nodea* we can point it at localhost because etcd is running on nodea.

```
laptop$ ssh -i k8s-adv-student.pem ubuntu@<external-ip>

...

ubuntu@nodea:~$ cd ~/goProjects/src/github.com/coreos/flannel

ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ sudo ./dist/flanneld --etcd-
endpoints=http://localhost:2379

I0331 01:44:17.321923    27251 main.go:516] Determining IP address of default interface
I0331 01:44:17.322139    27251 main.go:529] Using interface with name eth0 and address 172.31.28.198
I0331 01:44:17.322158    27251 main.go:546] Defaulting external address to interface address (172.31.28.198)
I0331 01:44:17.322220    27251 main.go:244] Created subnet manager: Etcd Local Manager with Previous Subnet: None
I0331 01:44:17.322234    27251 main.go:247] Installing signal handlers
I0331 01:44:17.323467    27251 main.go:388] Found network config - Backend type: aws-vpc
I0331 01:44:17.323506    27251 awsvpc.go:88] Backend configured as: %s{"Type": "aws-vpc"}
I0331 01:44:17.324533    27251 local_manager.go:234] Picking subnet in range 10.1.1.0 ... 10.1.255.0
I0331 01:44:17.325194    27251 local_manager.go:220] Allocated lease (10.1.40.0/24) to current node (172.31.28.198)
I0331 01:44:17.516755    27251 awsvpc.go:322] Found eni-0eb2bd30210caf566 that has 172.31.28.198 IP address.
W0331 01:44:17.564435    27251 awsvpc.go:134] failed to disable SourceDestCheck on eni-0eb2bd30210caf566:
UnauthorizedOperation: You are not authorized to perform this operation.
    status code: 403, request id: 99cf00a9-233d-4e0b-9b8f-3c35692d9d85.
I0331 01:44:17.564483    27251 awsvpc.go:79] Route table configured: false
I0331 01:44:17.724500    27251 awsvpc.go:141] Found route table rtb-6f23030a.
I0331 01:44:17.725009    27251 awsvpc.go:59] RouteTableID configured as string: %srtb-6f23030a
I0331 01:44:18.062726    27251 awsvpc.go:256] Route added to table rtb-6f23030a: 10.1.40.0/24 - eni-
0eb2bd30210caf566.
I0331 01:44:18.062759    27251 main.go:311] Changing default FORWARD chain policy to ACCEPT
I0331 01:44:18.062870    27251 main.go:319] Wrote subnet file to /run/flannel/subnet.env
I0331 01:44:18.062885    27251 main.go:323] Running backend.
```

```
I0331 01:44:18.063577 27251 main.go:431] Waiting for 22h59m59.261154778s to renew lease
I0331 01:44:18.065186 27251 iptables.go:145] Some iptables rules are missing; deleting and recreating rules
I0331 01:44:18.065204 27251 iptables.go:167] Deleting iptables rule: -s 10.1.0.0/16 -j ACCEPT
I0331 01:44:18.065919 27251 iptables.go:167] Deleting iptables rule: -d 10.1.0.0/16 -j ACCEPT
I0331 01:44:18.066800 27251 iptables.go:155] Adding iptables rule: -s 10.1.0.0/16 -j ACCEPT
I0331 01:44:18.068365 27251 iptables.go:155] Adding iptables rule: -d 10.1.0.0/16 -j ACCEPT
```

Take note of the subnet flannel uses (in the example: `Allocated lease, 10.1.40.0/24`), this is what we will configure the Docker daemon to use (later) instead of the default bridge subnet on *nodea*.

## Configure Docker

On *nodea*, remove the old bridge IP address. First shutdown docker:

```
ubuntu@nodea:~$ sudo systemctl stop docker

ubuntu@nodea:~$
```

Next, find the IP of the docker0 bridge and delete it:

```
ubuntu@nodea:~$ ip a show docker0

3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:1d:ce:ed:23 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:1dff:fece:ed23/64 scope link
        valid_lft forever preferred_lft forever

ubuntu@nodea:~$ sudo ip addr del 172.17.0.1/16 dev docker0

ubuntu@nodea:~$
```

When flannel starts on each node, it lays out a configuration file that others can use.

On *nodea*:

```
ubuntu@nodea:~$ cat /var/run/flannel/subnet.env

FLANNEL_NETWORK=10.1.0.0/16
FLANNEL_SUBNET=10.1.40.1/24
FLANNEL_MTU=9001
FLANNEL_IPMASQ=false
ubuntu@nodea:~$
```

Now create the Docker daemon configuration file on *nodea* so that docker assigns the docker0 bridge the value of FLANNEL\_SUBNET (in the example: 10.1.40.1/24) with the mtu of FLANNEL\_MTU (in the example: 9001):

```
ubuntu@nodea:~$ sudo vim /etc/docker/daemon.json

ubuntu@nodea:~$ cat /etc/docker/daemon.json

{
    "bip": "10.1.40.1/24",
    "mtu": 9001
}
ubuntu@nodea:~$
```

Restart docker:

```
ubuntu@nodea:~$ sudo systemctl start docker

ubuntu@nodea:~$ sudo systemctl status docker

● docker.service – Docker Application Container Engine
   Loaded: loaded (/lib/systemd/system/docker.service; enabled; vendor preset: enabled)
   Active: active (running) since Sun 2019-03-31 01:54:35 UTC; 7s ago
     Docs: https://docs.docker.com
   Main PID: 27545 (dockerd)
    Tasks: 10
   Memory: 31.6M
      CPU: 198ms
```

```
CGroup: /system.slice/docker.service
└─27545 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock
```

```
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.713566470Z" level=info msg="Graph migration to content-addressability took 0.00 seconds"
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.713727890Z" level=warning msg="Your kernel does not support swap memory limit"
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.713776454Z" level=warning msg="Your kernel does not support cgroup rt period"
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.713785804Z" level=warning msg="Your kernel does not support cgroup rt runtime"
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.714109539Z" level=info msg="Loading containers: start."
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.809428012Z" level=info msg="Loading containers: done."
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.836372066Z" level=info msg="Docker daemon" commit=774a1f4 graphdriver(s)=overlay2 version=18.09.3
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.836445676Z" level=info msg="Daemon has completed initialization"
Mar 31 01:54:35 nodea dockerd[27545]: time="2019-03-31T01:54:35.847890997Z" level=info msg="API listen on /var/run/docker.sock"
Mar 31 01:54:35 nodea systemd[1]: Started Docker Application Container Engine.
ubuntu@nodea:~$
```

Now run a test container, and inspect its eth0 interface:

```
ubuntu@nodea:~$ docker container run --rm ubuntu:14.04 ip a show eth0

Unable to find image 'ubuntu:14.04' locally
14.04: Pulling from library/ubuntu
e082d4499130: Pull complete
371450624c9e: Pull complete
c8a555b3a57c: Pull complete
1456d810d42e: Pull complete
Digest: sha256:6612de24437f6f01d6a2988ed9a36b3603df06e8d2c0493678f3ee696bc4bb2d
Status: Downloaded newer image for ubuntu:14.04
58: eth0@if59: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc noqueue state UP group default
    link/ether 02:42:0a:01:23:02 brd ff:ff:ff:ff:ff:ff
    inet 10.1.40.2/24 brd 10.1.40.255 scope global eth0
        valid_lft forever preferred_lft forever
ubuntu@nodea:~$
```



Great! Our container is connected to the new Flannel network!

## Copy flanneld to nodeb

Copy the flanneld executable to *nodeb*, this way we don't need to install all the compilation tools again. Follow these steps to copy the binary.

On *nodeb* run:

```
netcat -l -p 7000 | tar xv
```

On *nodea* run:

```
ubuntu@nodea:~$ cd ~/goProjects/src/github.com/coreos/flannel
ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ tar cf - dist/flanneld | nc nodeb 7000
```

## Start flanneld on nodeb

On *nodeb* start up the flanneld daemon **in its own terminal**. This time we will set the etcd endpoint to the hostname or IP of *nodea*:

```
ubuntu@nodeb:~$ sudo ./dist/flanneld --etcd-endpoints=http://nodea:2379

I0331 01:59:16.768732 14246 main.go:516] Determining IP address of default interface
I0331 01:59:16.769057 14246 main.go:529] Using interface with name eth0 and address 172.31.30.148
I0331 01:59:16.769078 14246 main.go:546] Defaulting external address to interface address (172.31.30.148)
I0331 01:59:16.769195 14246 main.go:244] Created subnet manager: Etcd Local Manager with Previous Subnet: None
I0331 01:59:16.769211 14246 main.go:247] Installing signal handlers
I0331 01:59:16.771005 14246 main.go:388] Found network config - Backend type: aws-vpc
I0331 01:59:16.771035 14246 awsvpc.go:88] Backend configured as: %s{"Type": "aws-vpc"}
I0331 01:59:16.772415 14246 local_manager.go:234] Picking subnet in range 10.1.1.0 ... 10.1.255.0
I0331 01:59:16.773135 14246 local_manager.go:220] Allocated lease (10.1.87.0/24) to current node (172.31.30.148)
I0331 01:59:16.991245 14246 awsvpc.go:322] Found eni-073753aa399a1de5f that has 172.31.30.148 IP address.
W0331 01:59:17.038623 14246 awsvpc.go:134] failed to disable SourceDestCheck on eni-073753aa399a1de5f:
UnauthorizedOperation: You are not authorized to perform this operation.
        status code: 403, request id: 2e5d6279-d54c-4af5-ab2a-45416d11403c.
I0331 01:59:17.038775 14246 awsvpc.go:79] Route table configured: false
```

```

I0331 01:59:17.211710 14246 awsvpc.go:141] Found route table rtb-6f23030a.
I0331 01:59:17.212613 14246 awsvpc.go:59] RouteTableID configured as string: %srtb-6f23030a
I0331 01:59:17.555961 14246 awsvpc.go:256] Route added to table rtb-6f23030a: 10.1.87.0/24 - eni-
073753aa399a1de5f.
I0331 01:59:17.555989 14246 main.go:311] Changing default FORWARD chain policy to ACCEPT
I0331 01:59:17.556105 14246 main.go:319] Wrote subnet file to /run/flannel/subnet.env
I0331 01:59:17.556121 14246 main.go:323] Running backend.
I0331 01:59:17.557071 14246 main.go:431] Waiting for 22h59m59.219057887s to renew lease
I0331 01:59:17.557978 14246 iptables.go:145] Some iptables rules are missing; deleting and recreating rules
I0331 01:59:17.557995 14246 iptables.go:167] Deleting iptables rule: -s 10.1.0.0/16 -j ACCEPT
I0331 01:59:17.558792 14246 iptables.go:167] Deleting iptables rule: -d 10.1.0.0/16 -j ACCEPT
I0331 01:59:17.559622 14246 iptables.go:155] Adding iptables rule: -s 10.1.0.0/16 -j ACCEPT
I0331 01:59:17.561129 14246 iptables.go:155] Adding iptables rule: -d 10.1.0.0/16 -j ACCEPT

```

Notice the subnet (Allocated lease), 10.1.87.0/24, this is what we will configure (later) Docker daemon to use versus the default bridge subnet on *nodeb* as well.

Now, with the *nodeb* leased subnet, follow the same procedure, changing the existing docker daemon configuration file:

```

ubuntu@nodeb:~$ sudo systemctl stop docker

ubuntu@nodeb:~$ ip a show docker0

3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:1d:ce:ed:23 brd ff:ff:ff:ff:ff:ff
    inet 172.18.0.1/16 brd 172.18.255.255 scope global docker0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:1dff:fece:ed23/64 scope link
        valid_lft forever preferred_lft forever

ubuntu@nodeb:~$ sudo ip addr del 172.18.0.1/16 dev docker0

ubuntu@nodeb:~$ cat /var/run/flannel/subnet.env

FLANNEL_NETWORK=10.1.0.0/16
FLANNEL_SUBNET=10.1.87.1/24
FLANNEL_MTU=9001
FLANNEL_IPMASQ=false
ubuntu@nodeb:~$

```

Now create the Docker daemon configuration file on *nodeb* so that docker assigns the docker0 bridge the value of FLANNEL\_SUBNET (in the example: 10.1.87.1/24) with the mtu of FLANNEL\_MTU (in the example: 9001):

```
ubuntu@nodeb:~$ sudo vim /etc/docker/daemon.json
ubuntu@nodeb:~$ sudo cat /etc/docker/daemon.json
```

```
{
    "bip": "10.1.87.1/24",
    "mtu": 9001
}
```

```
ubuntu@nodeb:~$ sudo systemctl start docker
```

```
ubuntu@nodeb:~$ sudo systemctl status docker
```

● **docker.service** – Docker Application Container Engine

Loaded: loaded (/lib/systemd/system/docker.service; enabled; vendor preset: enabled)

Active: active (running) since Sun 2019-03-31 02:01:46 UTC; 7s ago

Docs: <https://docs.docker.com>

Main PID: 14336 (dockerd)

Tasks: 10

Memory: 31.2M

CPU: 195ms

CGroup: /system.slice/docker.service

└─14336 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.289437235Z" level=info msg="Graph migration to content-addressability"
```

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.289605407Z" level=warning msg="Your kernel does not support swap memo"
```

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.289645515Z" level=warning msg="Your kernel does not support cgroup rt"
```

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.289655564Z" level=warning msg="Your kernel does not support cgroup rt"
```

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.290076894Z" level=info msg="Loading containers: start."
```

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.382212666Z" level=info msg="Loading containers: done."
```

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.418440682Z" level=info msg="Docker daemon"
commit=d14af54 graphdriver
```

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.418512793Z" level=info msg="Daemon has completed initialization"
```

```
Mar 31 02:01:46 nodeb dockerd[14336]: time="2019-03-31T02:01:46.430955768Z" level=info msg="API listen on /var/run/docker.sock"
Mar 31 02:01:46 nodeb systemd[1]: Started Docker Application Container Engine.

ubuntu@nodeb:~$ docker container run --rm ubuntu:14.04 ip a show eth0

Unable to find image 'ubuntu:14.04' locally
14.04: Pulling from library/ubuntu
e082d4499130: Pull complete
371450624c9e: Pull complete
c8a555b3a57c: Pull complete
1456d810d42e: Pull complete
Digest: sha256:6612de24437f6f01d6a2988ed9a36b3603df06e8d2c0493678f3ee696bc4bb2d
Status: Downloaded newer image for ubuntu:14.04
20: eth0@if21: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc noqueue state UP group default
    link/ether 02:42:0a:01:0d:02 brd ff:ff:ff:ff:ff:ff
    inet 10.1.87.2/24 brd 10.1.87.255 scope global eth0
        valid_lft forever preferred_lft forever
ubuntu@nodeb:~$
```

## 2. Investigate

### Investiate flannel Network

We saw earlier how flannel uses etcd to store information about its network. We are using many default settings, but all of those stem from the single network cidr we entered earlier. Now that flannel is up and running on both nodes, lets look at the networking contstructs.

On nodea, review the interfaces:

```
ubuntu@nodea:~/goProjects/src/github.com/coreos/flannel$ cd ~

ubuntu@nodea:~$ ip a show

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
```

```
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc mq state UP group default qlen 1000
    link/ether 02:ef:63:d5:3b:be brd ff:ff:ff:ff:ff:ff
    inet 172.31.28.198/20 brd 172.31.31.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::ef:63ff:fed5:3bbe/64 scope link
        valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:4f:8a:e5:22 brd ff:ff:ff:ff:ff:ff
    inet 10.1.40.1/24 brd 10.1.40.255 scope global docker0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:4fff:fe8a:e522/64 scope link
        valid_lft forever preferred_lft forever
ubuntu@nodea:~$
```

review routing information:

```
ubuntu@nodea:~$ ip route show

default via 172.31.16.1 dev eth0
10.1.40.0/24 dev docker0 proto kernel scope link src 10.1.40.1 linkdown
172.18.0.0/16 via 172.31.30.148 dev eth0
172.31.16.0/20 dev eth0 proto kernel scope link src 172.31.28.198
ubuntu@nodea:~$
```

Hmmm, not much information here. Do the same on nodeb and take note of the IP information.

```
ubuntu@nodeb:~$ ip a show

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc mq state UP group default qlen 1000
    link/ether 02:d8:c0:66:a6:b8 brd ff:ff:ff:ff:ff:ff
    inet 172.31.30.148/20 brd 172.31.31.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::d8:c0ff:fe66:a6b8/64 scope link
```

```
valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
link/ether 02:42:5f:57:41:79 brd ff:ff:ff:ff:ff:ff
inet 10.1.87.1/24 brd 10.1.87.255 scope global docker0
    valid_lft forever preferred_lft forever
inet6 fe80::42:5fff:fe57:4179/64 scope link
    valid_lft forever preferred_lft forever
```

```
ubuntu@nodeb:~$ ip route show
```

```
default via 172.31.16.1 dev eth0
10.1.87.0/24 dev docker0 proto kernel scope link src 10.1.87.1 linkdown
172.17.0.0/16 via 172.31.28.198 dev eth0
172.31.16.0/20 dev eth0 proto kernel scope link src 172.31.30.148
ubuntu@nodeb:~$
```

In AWS, flannel can be configured to work with Amazon VPC. When we set network details in etcd, we told it to use the "type" "aws-vpc", so the routing table is stored there, for example:

rtb-a61df4cf

Summary

Routes

Subnet Associations

Route Propagation

Tags

Edit

View:

All rules



Destination	Target	Status	Propagated
172.31.0.0/16	local	Active	No
0.0.0.0/0	igw-e80ae581	Active	No
10.1.40.0/24	eni-0ae8131d52a4421ff / i-05a3d3717460b550e	Active	No
10.1.87.0/24	eni-043353e0bbb8ffd26 / i-04d31c712e7c8dfa5	Active	No

This requires an AWS IAM policy to be applied to all of the instances:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "ec2:CreateRoute",
        "ec2>DeleteRoute",
        "ec2:ReplaceRoute"
      ],
      "Resource": [
```

```

        "*"
    ],
    {
        "Effect": "Allow",
        "Action": [
            "ec2:DescribeRouteTables",
            "ec2:DescribeInstances"
        ],
        "Resource": "*"
    }
]
}

```

This step was taken care of by the instructor.

To test that we are being routed correctly, try pinging the docker0 gateway between the nodes.

From nodea to nodeb.

```

ubuntu@nodea:~$ ping -c 1 10.1.87.1

PING 10.1.87.1 (10.1.87.1) 56(84) bytes of data.
64 bytes from 10.1.87.1: icmp_seq=1 ttl=64 time=0.363 ms

--- 10.1.87.1 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.363/0.363/0.363/0.000 ms
ubuntu@nodea:~$

```

and reverse

```

ubuntu@nodeb:~$ ping -c 1 10.1.40.1

PING 10.1.40.1 (10.1.40.1) 56(84) bytes of data.
64 bytes from 10.1.40.1: icmp_seq=1 ttl=64 time=0.404 ms

--- 10.1.40.1 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms

```



```
rtt min/avg/max/mdev = 0.404/0.404/0.404/0.000 ms
ubuntu@nodea:~$
```

Now that we have routing working, lets try with containers.

## Testing With Docker

Overlay networks are used to augment traditional Docker networking, which belong to a single machine a-la docker0 bridge. Our goal is to contact the remote Docker container from another node.

On nodeb run the following command:

```
ubuntu@nodeb:~$ docker container run --rm -it -p 8889 ubuntu:14.04 nc -l 0.0.0.0 8889

...
```

In another terminal, get your container's IP address:

```
ubuntu@nodeb:~$ ID=$(docker container ls -q) && echo $ID

ff14668548fe

ubuntu@nodeb:~$ docker container inspect -f '{{.NetworkSettings.Networks.bridge.IPAddress }}' $ID

10.1.87.2
ubuntu@nodeb:~$
```

On nodea run the following command using the IP address returned by the inspect command above:



```
ubuntu@nodea:~$ telnet 10.1.87.2 8889
Trying 10.1.87.2...
Connected to 10.1.87.2.
Escape character is '^]'.
```

```
hello!
```

```
ubuntu@nodeb:~$ docker container run --rm -it -p 8889 ubuntu:14.04 nc -l 0.0.0.0 8889
```

```
hello!
```

Try typing on either node's terminal. You have a simple chat application via flannel with Docker.

To exit this example, hit ctrl+] (  ) on nodea followed by  at the telnet prompt:

```
ubuntu@nodea:~$ telnet 10.1.87.2 8889
Trying 10.1.87.2...
Connected to 10.1.87.2.
Escape character is '^]'.
hello!
hi
^]
telnet> q
Connection closed.
ubuntu@nodea:~$
```

Now lets run a website on nodeb.

```
ubuntu@nodeb:~$ docker container run -p 80 nginx:1.9.1
```

```
...
```

**In another terminal**, get your container's IP address:

```
ubuntu@nodeb:~$ ID=$(docker container ls -q) && echo $ID

eff4c3ae4594
```

```
ubuntu@nodeb:~$ docker container inspect -f '{{.NetworkSettings.Networks.bridge.IPAddress }}' $ID  
  
10.1.87.2  
ubuntu@nodeb:~$
```

Now on nodea:

```
ubuntu@nodea:~$ docker container run --rm -it centos  
  
[root@5526a8acdf9b /]# curl --head 10.1.87.2  
HTTP/1.1 200 OK  
Server: nginx/1.9.1  
Date: Thu, 03 May 2018 05:09:04 GMT  
Content-Type: text/html  
Content-Length: 612  
Last-Modified: Tue, 26 May 2015 15:02:09 GMT  
Connection: keep-alive  
ETag: "55648af1-264"  
Accept-Ranges: bytes  
  
[root@5526a8acdf9b /]
```

In your original terminal, you will see the request:

```
ubuntu@nodeb:~$ docker container run -p 80 nginx:1.9.1  
  
172.31.7.235 - - [03/May/2018:05:08:44 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.29.0" "-"
```

Success! Stop both containers.

## Testing With kubelet

Moving further up the chain, lets make sure this works with Kubernetes. We will use kubelet on nodea to launch a container, and in turn use nodeb to access it.

```
ubuntu@nodea:~$ vim poda.yaml
```

```
ubuntu@nodea:~$ cat poda.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  nodeName: nodea
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
      hostPort: 80
ubuntu@nodea:~$
```

Launch on nodea with kubelet directly:

```
ubuntu@nodea:~$ sudo $HOME/k8s/_output/bin/kubelet \
--allow-privileged=true \
--runtime-cgroups=/systemd/machine.slice \
--kubelet-cgroups=/systemd/machine.slice \
--pod-infra-container-image=k8s.gcr.io/pause:3.1 \
--pod-manifest-path=$HOME/poda.yaml

...
```

Find the nginx container's IP (remember it is assigned to the "pause" container!):

```
ubuntu@nodea:~$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
accf2668818d	nginx	"nginx -g 'daemon of...'"	5 minutes ago	Up 5 minutes	
k8s_nginx_nginx-nodea_default_349055bae6506cdf584cb0ae1cefa4d7_0					
0fd5138923c4	k8s.gcr.io/pause:3.1	"/pause"	5 minutes ago	Up 5 minutes	
0.0.0.0:80->80/tcp	k8s_POD_nginx-nodea_default_349055bae6506cdf584cb0ae1cefa4d7_0				

```
ubuntu@nodea:~$ CIP=$(docker container inspect $(docker container ls -q --filter=ancestor=k8s.gcr.io/pause:3.1) \
-f '{{.NetworkSettings.IPAddress}}') && echo $CIP
```

```
10.1.40.2
ubuntu@nodea:~$
```

On *nodeb*, lets access via curl.

```
ubuntu@nodeb:~$ curl --head $CIP

HTTP/1.1 200 OK
Server: nginx/1.15.10
Date: Sun, 31 Mar 2019 02:21:35 GMT
Content-Type: text/html
Content-Length: 612
Last-Modified: Tue, 26 Mar 2019 14:04:38 GMT
Connection: keep-alive
ETag: "5c9a3176-264"
Accept-Ranges: bytes
ubuntu@nodeb:~$
```

We've successfully added flannel networking! Be sure to remove your on-cluster resources.

Congratulations you have successfully completed the lab!

At this point you can restart kubelet (with the proper flags) and all of the master and node services, using flannel to provide networking to your cluster! Below are the commands for the rest of the services on nodea and nodeb.

nodea:

## api-server

```
sudo $HOME/k8s/_output/bin/kube-apiserver \
--etcd-servers=http://localhost:2379 \
--service-cluster-ip-range=10.0.0.0/16 \
--insecure-bind-address=0.0.0.0 \
--disable-admission-plugins=ServiceAccount
```

---

## controller-manager

```
$HOME/k8s/_output/bin/kube-controller-manager --kubeconfig=nodea.conf
```

## scheduler

```
$HOME/k8s/_output/bin/kube-scheduler --kubeconfig=nodea.conf
```

## kubelet on nodea

```
sudo rm -rf /var/lib/kubelet

sudo $HOME/k8s/_output/bin/kubelet \
--kubeconfig=nodea.conf \
--config=nodea.yaml \
--allow-privileged=true \
--runtime-cgroups=/systemd/machine.slice \
--kubelet-cgroups=/systemd/machine.slice \
--pod-infra-container-image=k8s.gcr.io/pause:3.1
```

## kube-proxy on nodea

```
sudo $HOME/k8s/_output/bin/kube-proxy --config=kube-proxy-config
```

nodeb:

## kubelet on nodeb

```
sudo rm -rf /var/lib/kubelet

sudo $HOME/kube-bin/kubelet \
--kubeconfig=nodeb.conf \
--config=nodeb.yaml \
--allow-privileged=true \
--runtime-cgroups=/systemd/machine.slice \
--kubelet-cgroups=/systemd/machine.slice \
--pod-infra-container-image=k8s.gcr.io/pause:3.1
```

## kube-proxy on nodeb

```
sudo $HOME/kube-bin/kube-proxy --config=kube-proxy-config
```

*Copyright (c) 2013-2019 RX-M LLC, Cloud Native Consulting, all rights reserved*